

**UCLA CS 118 Winter 2021**

Instructor: Giovanni Pau  
TAs:  
Hunter Dellaverson  
Eric Newberry

This chapter slide deck draws from different sources 7<sup>th</sup> and 8<sup>th</sup> edition of the textbook

Transport Layer: 3-1

1

## Chapter 3 Transport Layer

A note on the use of these PowerPoint slides:  
We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a lot of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR  
© All rights reserved. 1996-2016  
J.F. Kurose and K.W. Ross, All Rights Reserved

**Computer Networking: A Top Down Approach**  
7<sup>th</sup> edition  
Jim Kurose, Keith Ross  
Pearson/Addison Wesley  
April 2016  
Transport Layer

2-2

2

<https://powcoder.com>

## Chapter 3 Transport Layer

A note on the use of these PowerPoint slides:  
We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a lot of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

For a revision history, see the slide note for this page.

Thanks and enjoy! JFK/KWR  
All material copyright 1996-2020  
J.F. Kurose and K.W. Ross, All Rights Reserved

Transport Layer: 3-3

3

## Add WeChat powcoder

### Transport layer: overview

*Our goal:*

- understand principles behind transport layer services:
  - multiplexing, demultiplexing
  - reliable data transfer
  - flow control
  - congestion control
- learn about Internet transport layer protocols:
  - UDP: connectionless transport
  - TCP: connection-oriented reliable transport
  - TCP congestion control

Transport Layer: 3-4

4

## Transport layer: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer functionality

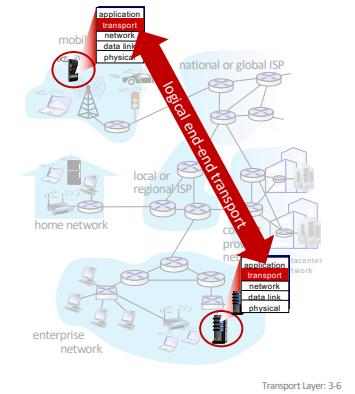


Transport Layer: 3-5

5

## Transport services and protocols

- provide *logical communication* between application processes running on different hosts
- transport protocols actions in end systems:
  - sender: breaks application messages into *segments*, passes to network layer
  - receiver: reassembles segments into messages, passes to application layer
- two transport protocols available to Internet applications
  - TCP, UDP



6

<https://powcoder.com>

## Transport vs. network layer services and protocols



### household analogy:

- 12 kids in Ann's house sending letters to 12 kids in Bill's house:
- hosts = houses
  - processes = kids
  - app messages = letters in envelopes

Transport Layer: 3-7

7

## Add WeChat powcoder

## Transport vs. network layer services and protocols

▪ **network layer:** logical communication between *hosts*

▪ **transport layer:** logical communication between *processes*

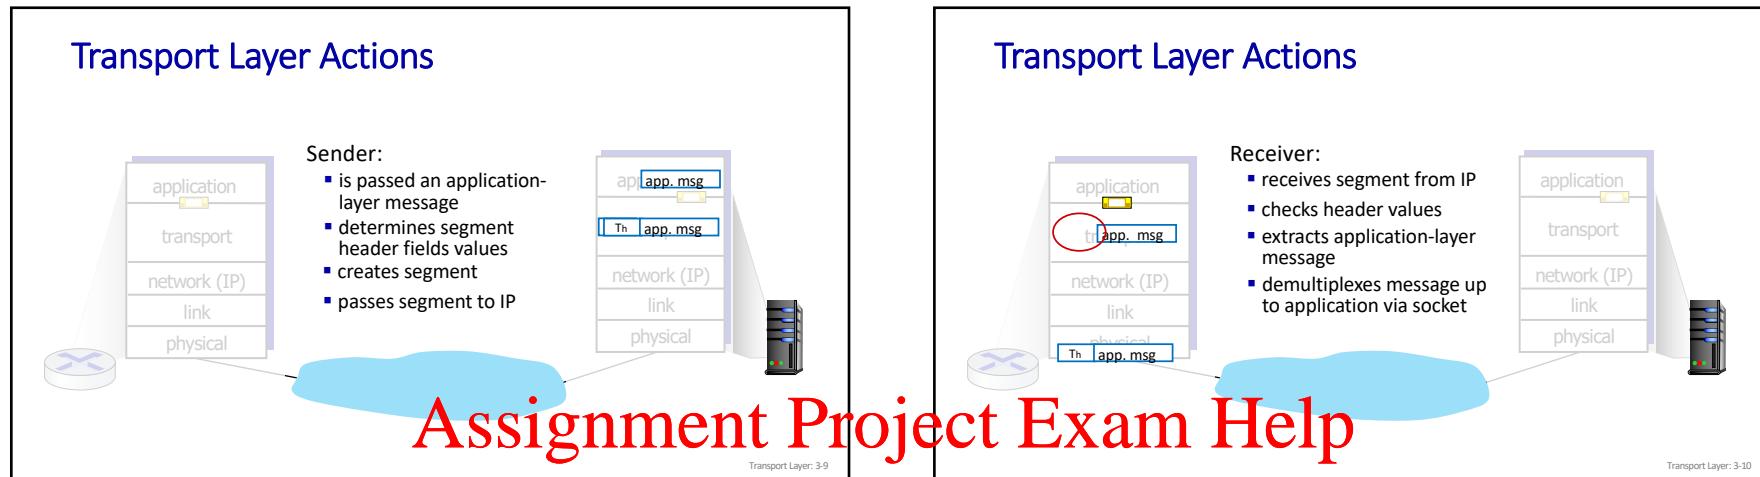
- relies on, enhances, network layer services

### household analogy:

- 12 kids in Ann's house sending letters to 12 kids in Bill's house:
- hosts = houses
  - processes = kids
  - app messages = letters in envelopes

Transport Layer: 3-8

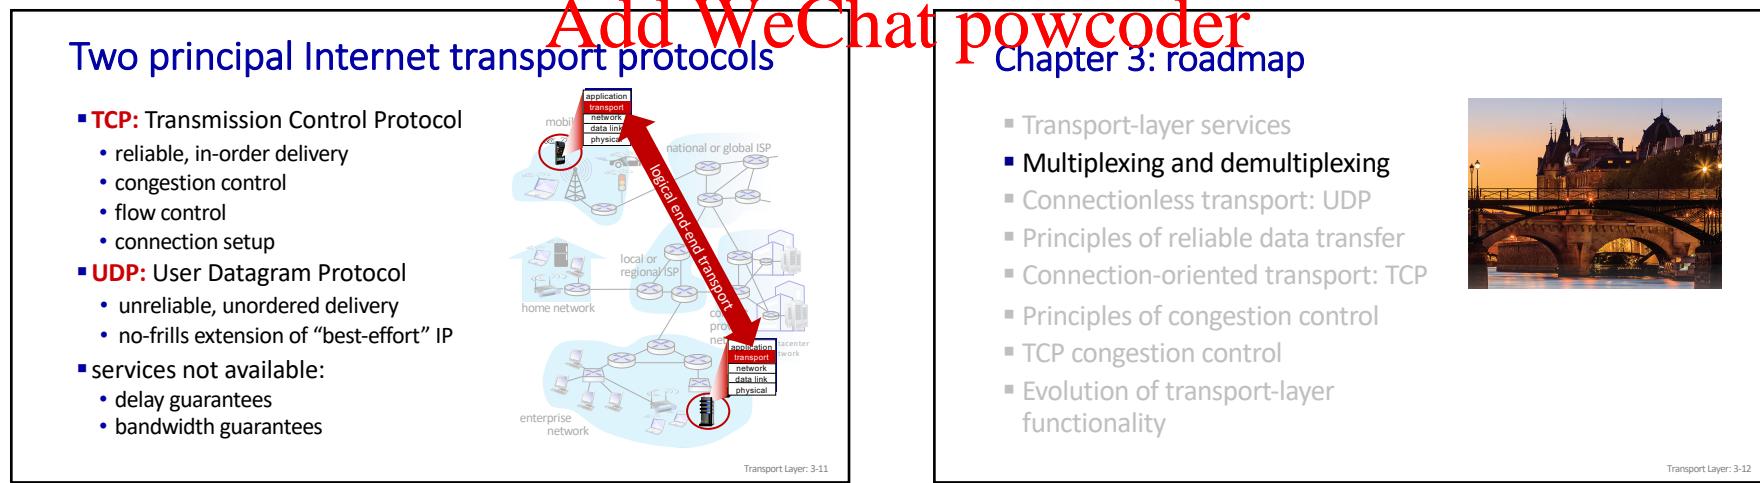
8



9

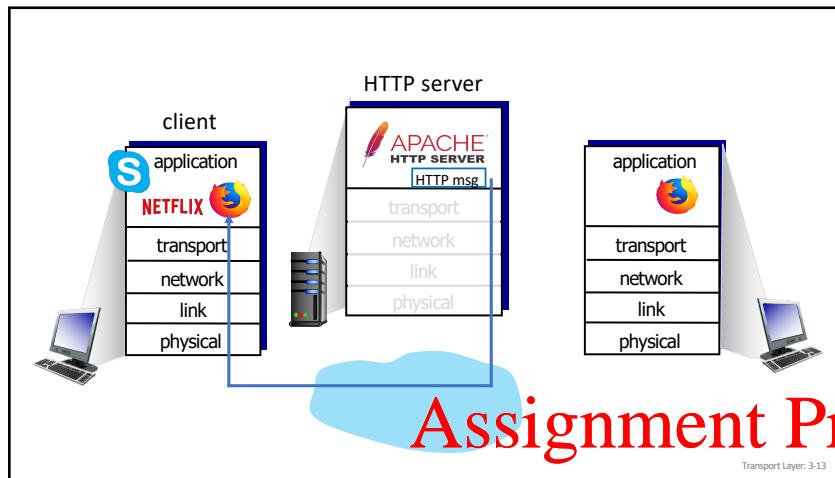
10

<https://powcoder.com>

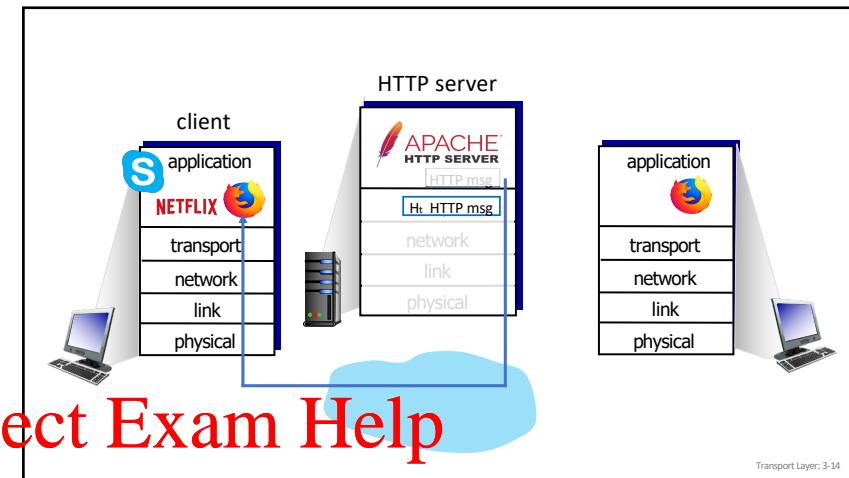


11

12



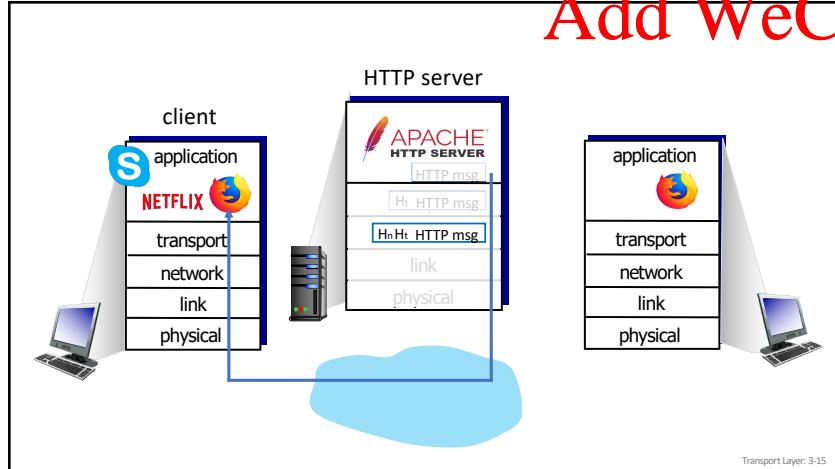
13



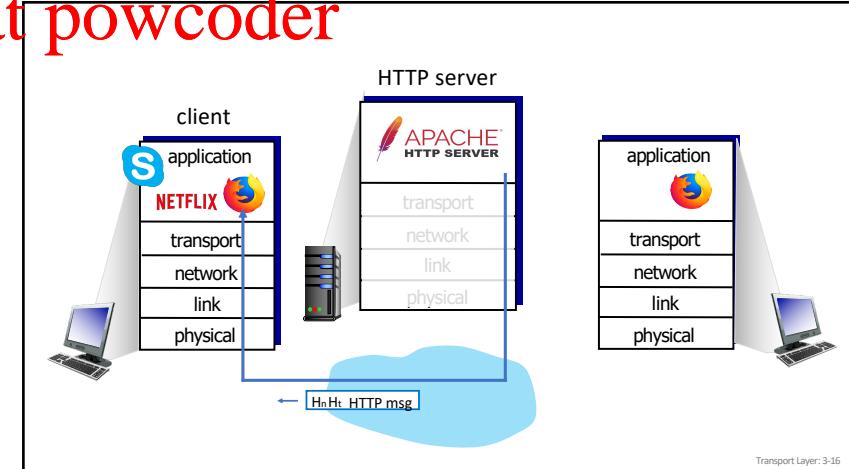
14

<https://powcoder.com>

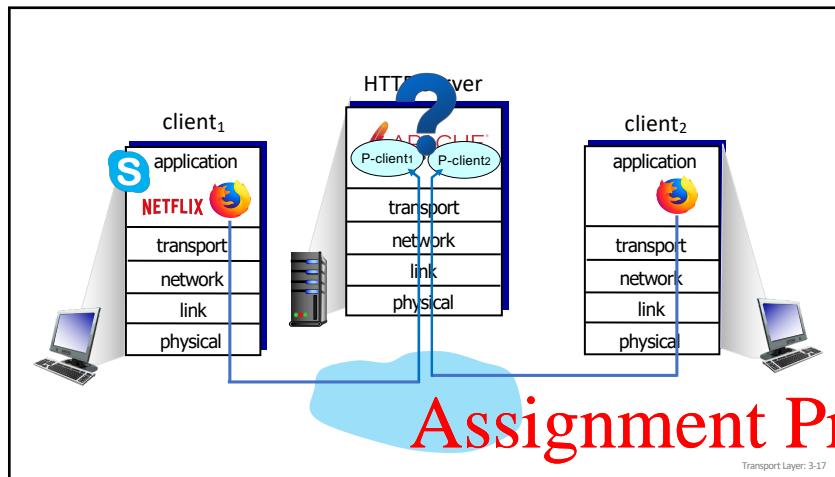
Add WeChat powcoder



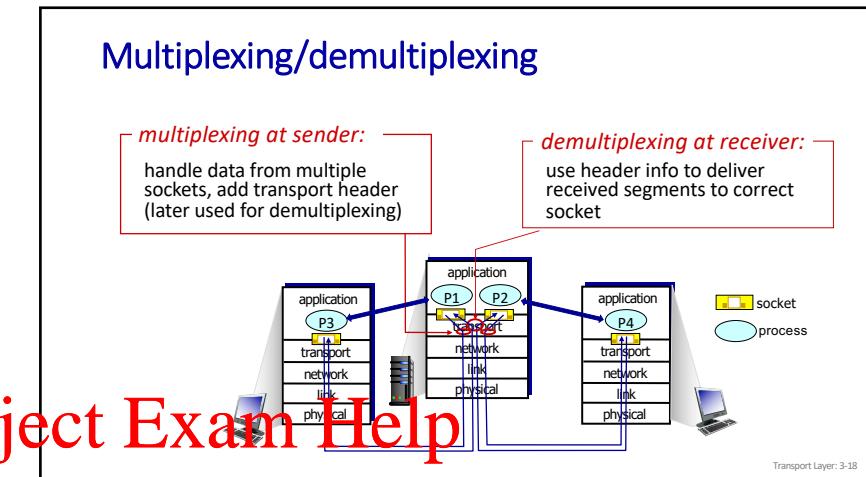
15



16

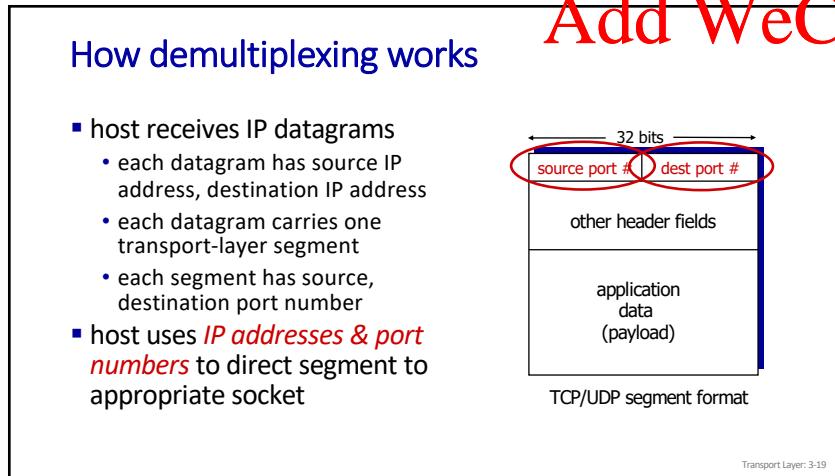


17

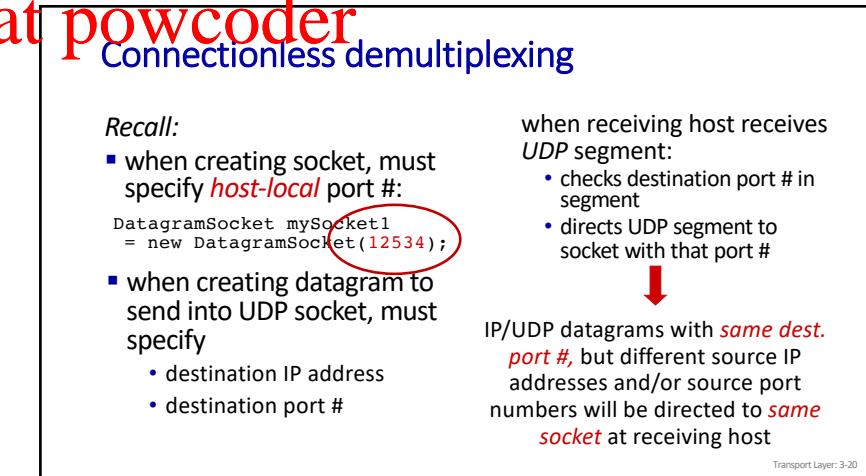


18

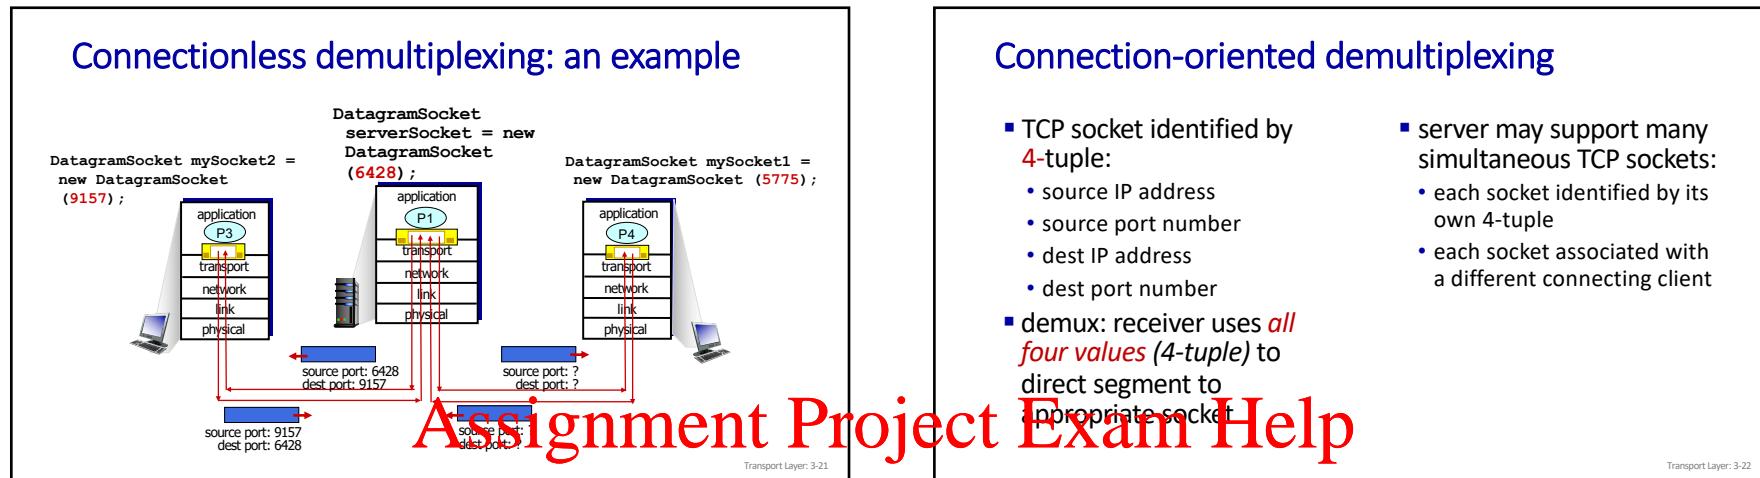
<https://powcoder.com>



19



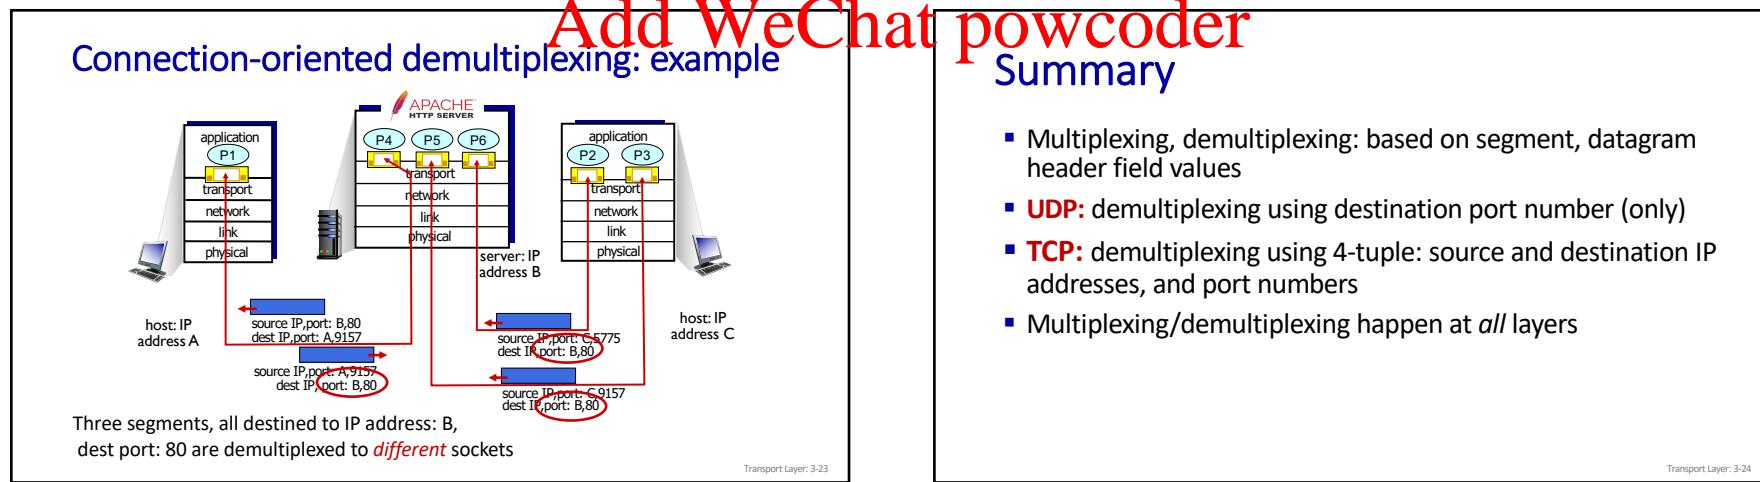
20



21

22

<https://powcoder.com>



23

24

## Connection-oriented demultiplexing

- TCP socket identified by **4-tuple**:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- demux: receiver uses *all four values* (4-tuple) to direct segment to appropriate socket

- server may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
  - each socket associated with a different connecting client

- ## Add WeChat powcoder Summary
- Multiplexing, demultiplexing: based on segment, datagram header field values
  - **UDP**: demultiplexing using destination port number (only)
  - **TCP**: demultiplexing using 4-tuple: source and destination IP addresses, and port numbers
  - Multiplexing/demultiplexing happen at *all* layers

### Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer functionality



Transport Layer: 3-25

# Assignment Project Exam Help

25

26

### UDP: User Datagram Protocol

- “no frills,” “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
  - lost
  - delivered out-of-order to app
- ***connectionless***:
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others

**Why is there a UDP?**

- no connection establishment (which can add RTT delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control
  - UDP can blast away as fast as desired!
  - can function in the face of congestion

### UDP: User Datagram Protocol

- UDP use:
  - streaming multimedia apps (loss tolerant, rate sensitive)
  - DNS
  - SNMP
  - HTTP/3
- if reliable transfer needed over UDP (e.g., HTTP/3):
  - add needed reliability at application layer
  - add congestion control at application layer

Add WeChat powcoder

27

28

### UDP: User Datagram Protocol [RFC 768]

INTERNET STANDARD  
RFC 768 J. Postel ISI 28 August 1980

User Datagram Protocol

Introduction

This User Datagram Protocol (UDP) is defined to make available a datagram mode of packet-switched computer communication in the environment of an interconnected set of computer networks. This protocol assumes that the Internet Protocol (IP) [1] is used as the underlying protocol.

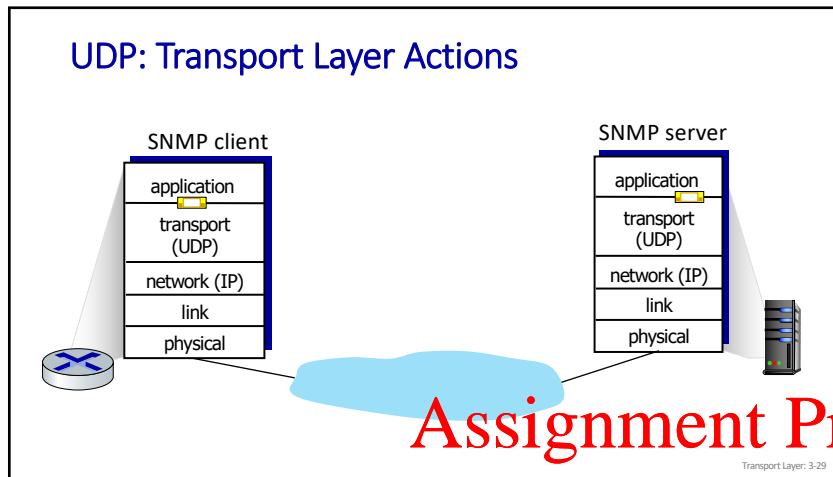
This protocol provides a procedure for application programs to send messages to other programs with a minimum of protocol mechanism. The protocols is transaction oriented, and delivery and duplicate protection are not guaranteed. Applications requiring reliable delivery of streams of data should use the Transmission Control Protocol (TCP) [2].

Format

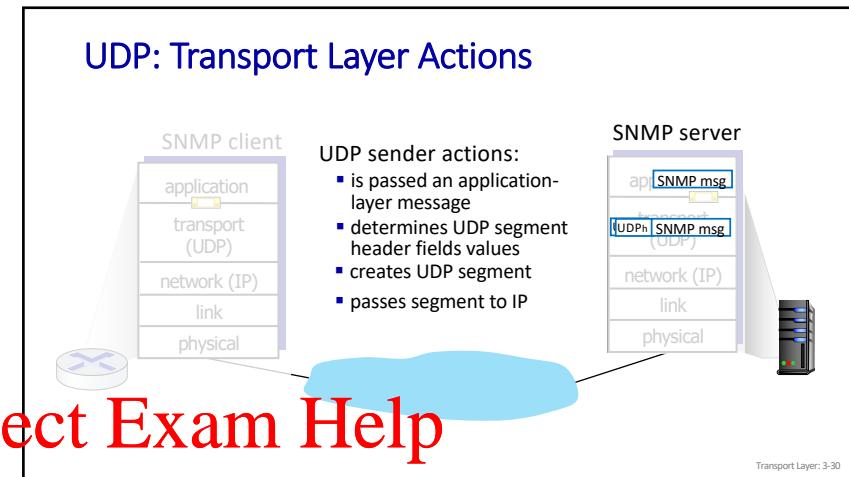
0	7 8	15 16	23 24	31
Source Port		Destination Port		
Length		Checksum		
data octets ...				

Transport Layer: 3-27

Transport Layer: 3-28

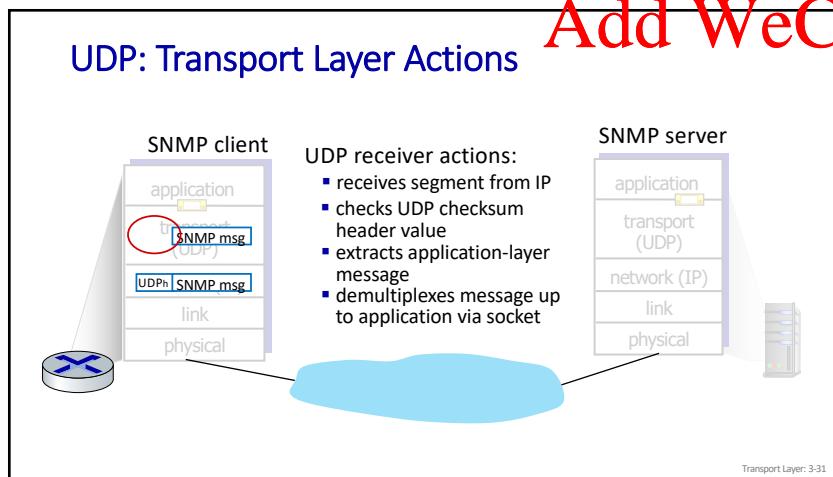


29

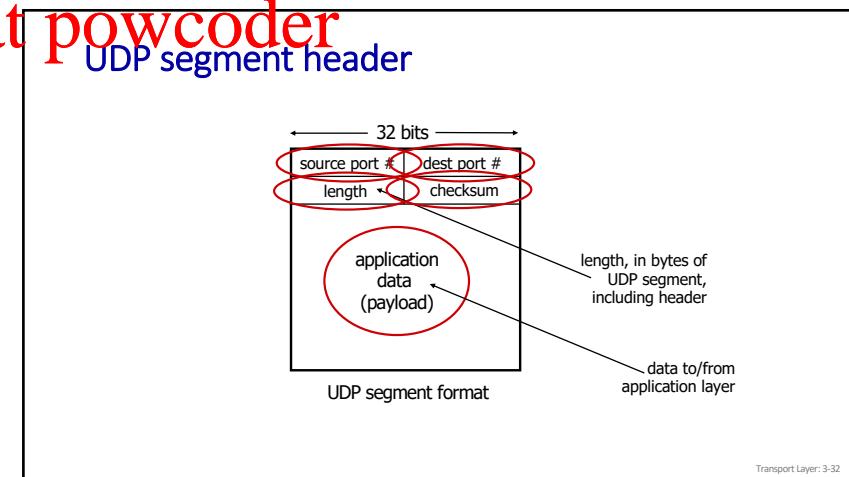


30

<https://powcoder.com>



31



## UDP checksum

**Goal:** detect errors (i.e., flipped bits) in transmitted segment

	1 <sup>st</sup> number	2 <sup>nd</sup> number	sum
Transmitted:	5	6	11
Received:	4	6	11

receiver-computed checksum ≠ sender-computed checksum (as received)

Assignment Project Exam Help

Transport Layer: 3-33

33

## UDP checksum

**Goal:** detect errors (i.e., flipped bits) in transmitted segment

**sender:**

- treat contents of UDP segment (including UDP header fields and IP addresses) as sequence of 16-bit integers
- checksum:** addition (one's complement sum) of segment content
- checksum value put into UDP checksum field

**receiver:**

- compute checksum of received segment
- check if computed checksum equals checksum field value:
  - Not equal - error detected
  - Equal - no error detected. *But maybe errors nonetheless? More later ....*

Transport Layer: 3-34

34

Assignment Project Exam Help

<https://powcoder.com>

## Internet checksum: an example

example: add two 16-bit integers

1 1 1 0 0 1 1 0 0 1 1 0 0 1 1 0	1 1 0 1 0 1 0 1 0 1 0 1 0 1 0 1
<hr/>	
wraparound ① 1 0 1 1 1 0 1 1 1 0 1 1 1 0 1 1	0 1
<hr/>	
sum 1 0 1 1 1 0 1 1 1 0 1 1 1 1 1 0 0	1 0
<hr/>	
checksum 0 1 0 0 0 1 0 0 0 1 0 0 0 0 1 1	0 1

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

\* Check out the online interactive exercises for more examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)

Transport Layer: 3-35

35

## Internet checksum: weak protection!

example: add two 16-bit integers

1 1 1 0 0 1 1 0 0 1 1 0 0 1 1 0	1 1 0 1 0 1 0 1 0 1 0 1 0 1 0 1
<hr/>	
wraparound ① 1 0 1 1 1 0 1 1 1 0 1 1 1 0 1 1	0 1 1 0 0 1
<hr/>	
sum 1 0 1 1 1 0 1 1 1 0 1 1 1 1 1 0 0	1 0
<hr/>	
checksum 0 1 0 0 0 1 0 0 0 1 0 0 0 0 1 1	0 1

Even though numbers have changed (bit flips), **no** change in checksum!

Transport Layer: 3-36

36

## Summary: UDP

- “no frills” protocol:
  - segments may be lost, delivered out of order
  - best effort service: “send and hope for the best”
- UDP has its pluses:
  - no setup/handshaking needed (no RTT incurred)
  - can function when network service is compromised
  - helps with reliability (checksum)
- build additional functionality on top of UDP in application layer (e.g., HTTP/3)

37

# Assignment Project Exam Help

## Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer services

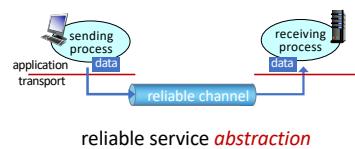


Transport Layer: 3-38

38

<https://powcoder.com>

## Principles of reliable data transfer

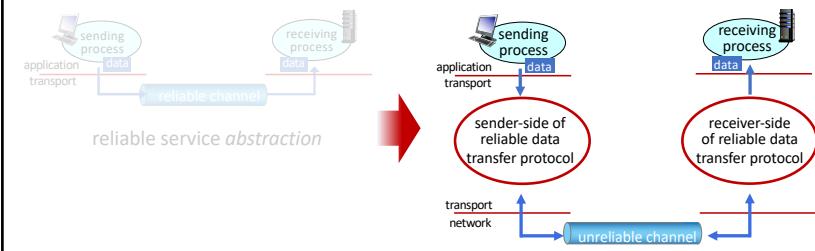


Transport Layer: 3-39

39

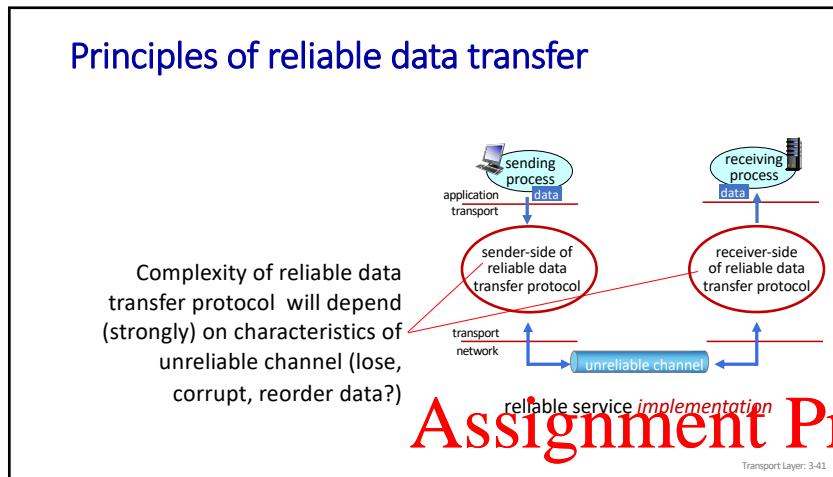
## Add WeChat powcoder

## Principles of reliable data transfer

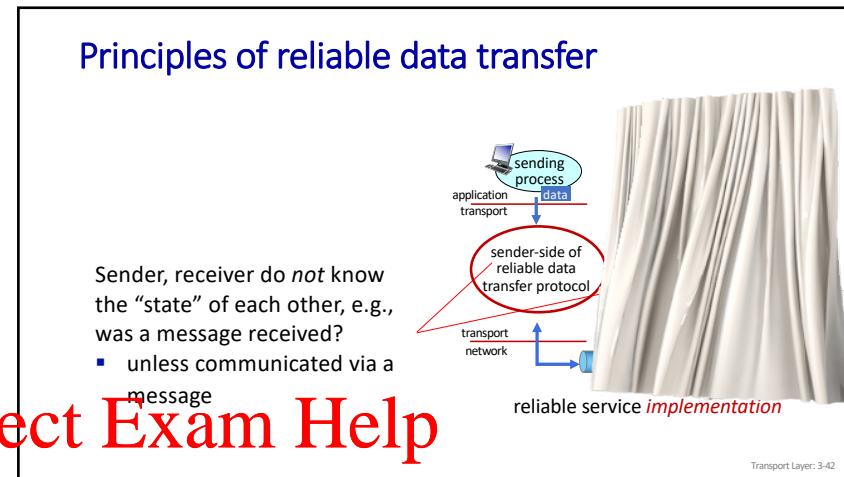


Transport Layer: 3-40

40

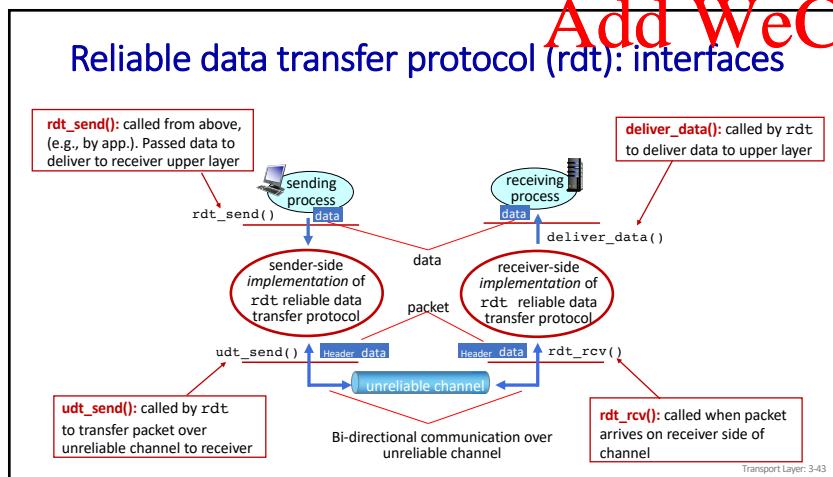


41

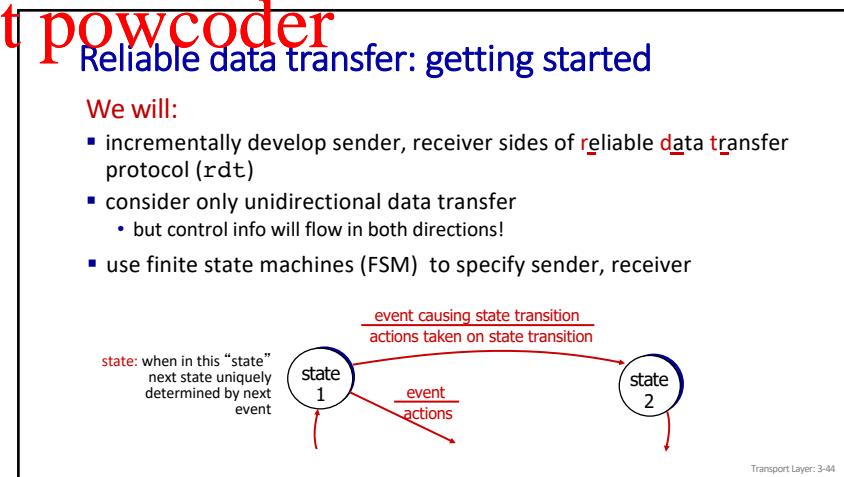


42

<https://powcoder.com>



43



44

**rdt1.0: reliable transfer over a reliable channel**

- underlying channel perfectly reliable
  - no bit errors
  - no loss of packets
- separate** FSMS for sender, receiver:
  - sender sends data into underlying channel
  - receiver reads data from underlying channel

**Assignment Project Exam Help**

Transport Layer: 3-45

**rdt2.0: channel with bit errors**

- underlying channel may flip bits in packet
  - checksum (e.g., Internet checksum) to detect bit errors
- the question:** how to recover from errors?

*How do humans recover from “errors” during conversation?*

Transport Layer: 3-46

45

46

<https://powcoder.com>

**rdt2.0: channel with bit errors**

- underlying channel may flip bits in packet
  - checksum to detect bit errors
- the question:** how to recover from errors?
  - acknowledgements (ACKs):** receiver explicitly tells sender that pkt received OK
  - negative acknowledgements (NAKs):** receiver explicitly tells sender that pkt had errors
  - sender **retransmits** pkt on receipt of NAK

**stop and wait**

sender sends one packet, then waits for receiver response

Transport Layer: 3-47

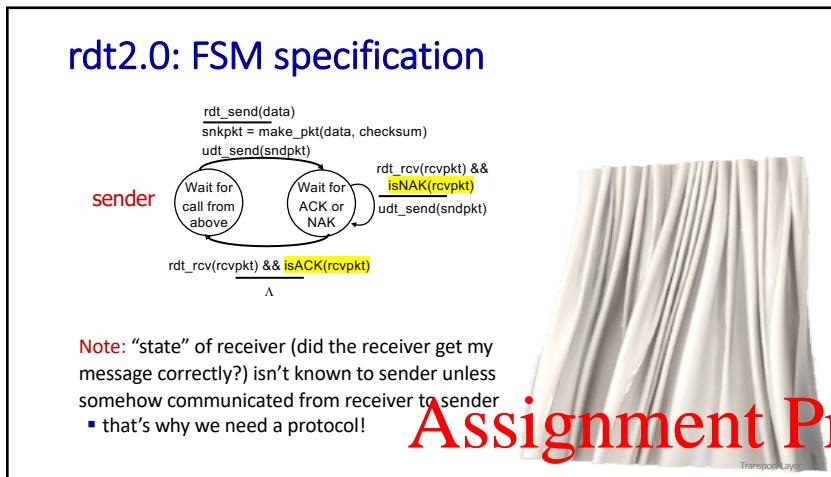
**Add WeChat powcoder**

**rdt2.0: FSM specifications**

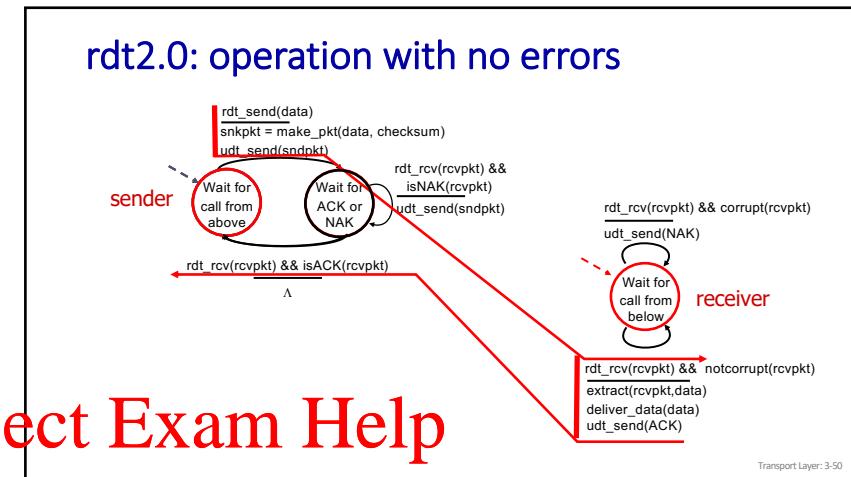
Transport Layer: 3-48

47

48

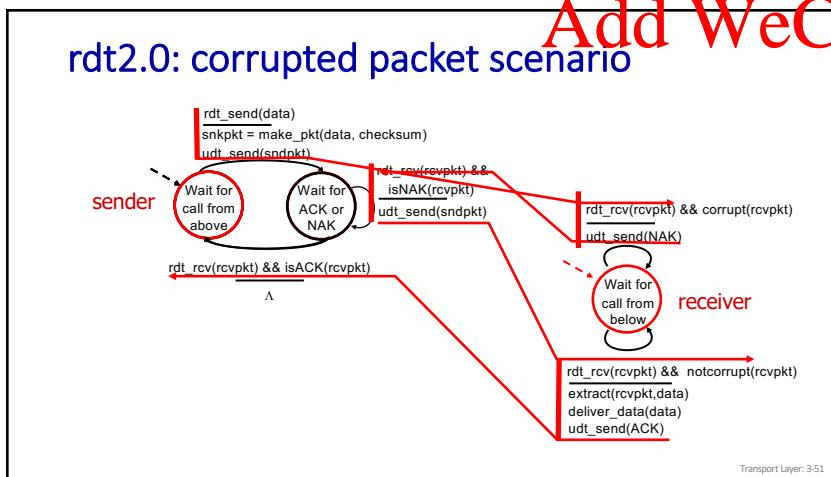


49

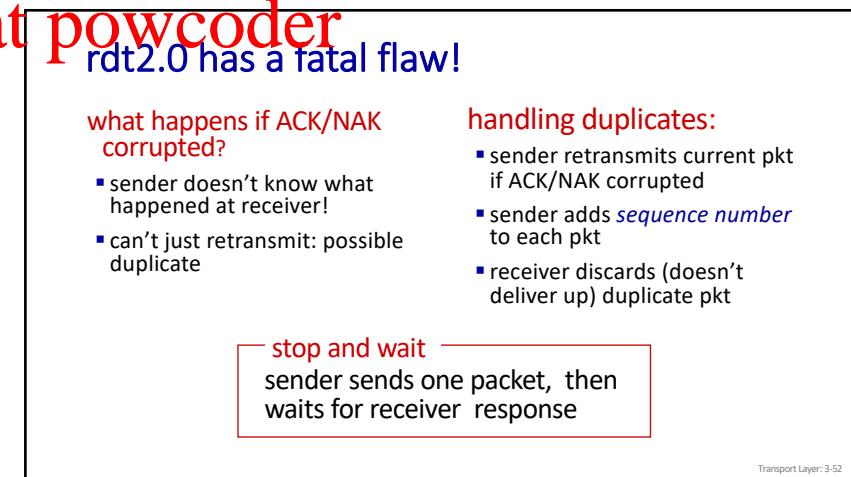


50

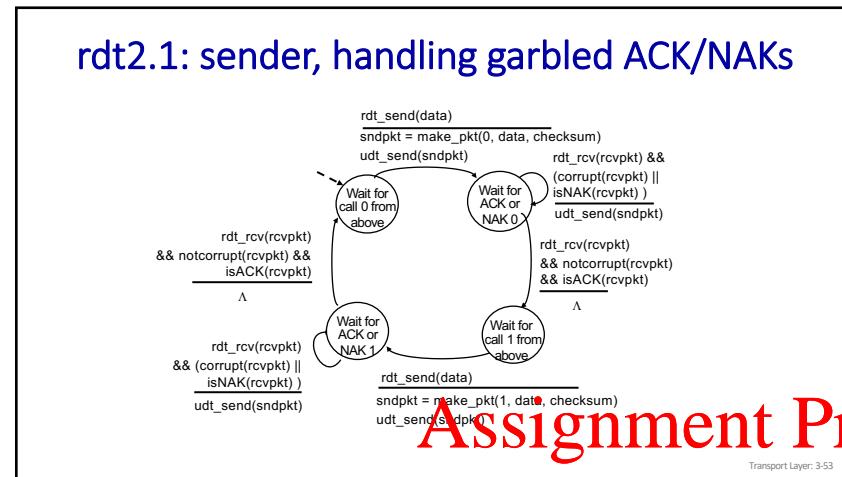
# Assignment Project Exam Help



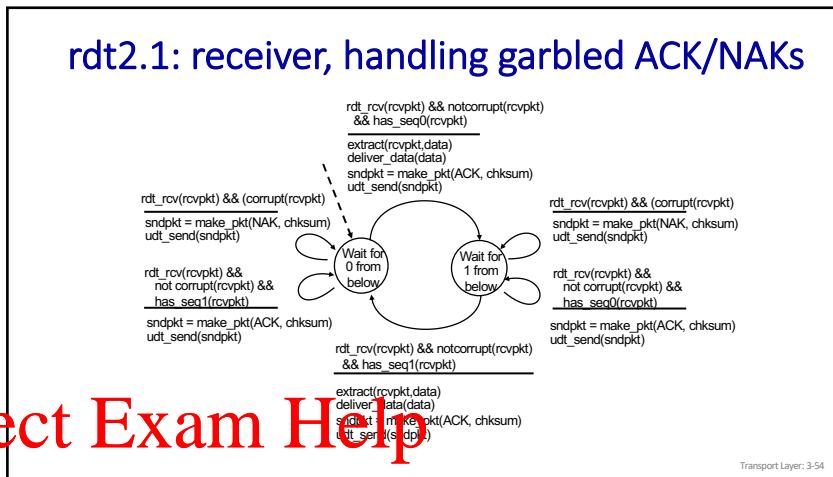
51



52



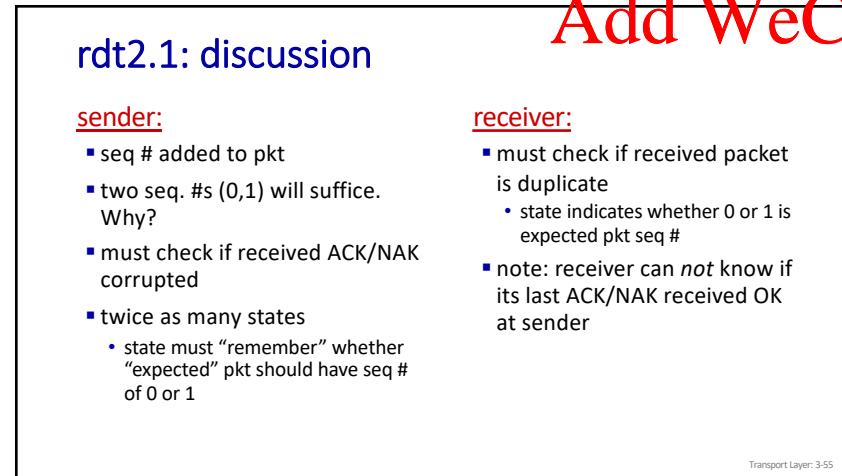
53



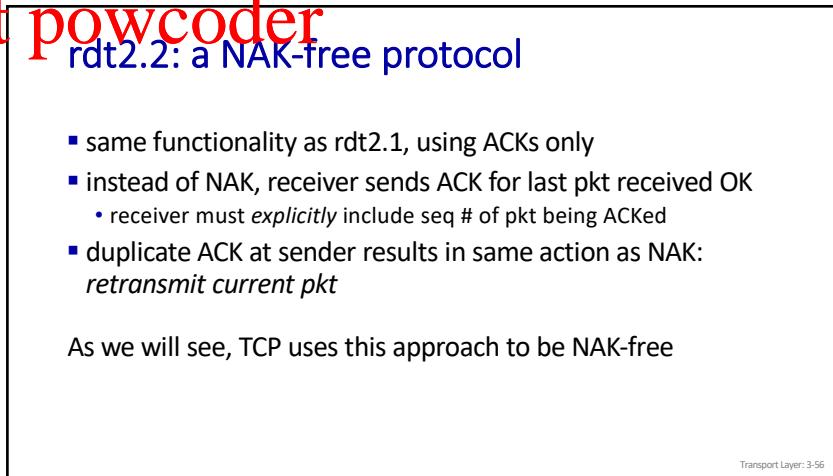
54

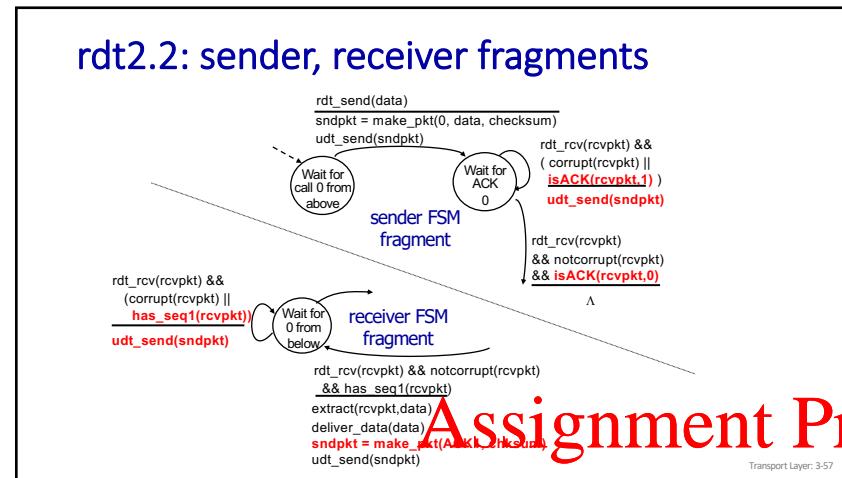
# Assignment Project Exam Help

<https://powcoder.com>

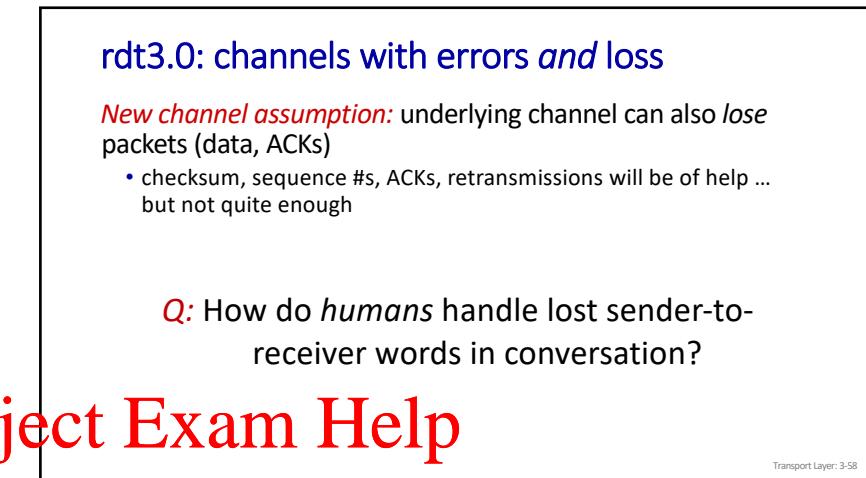


55



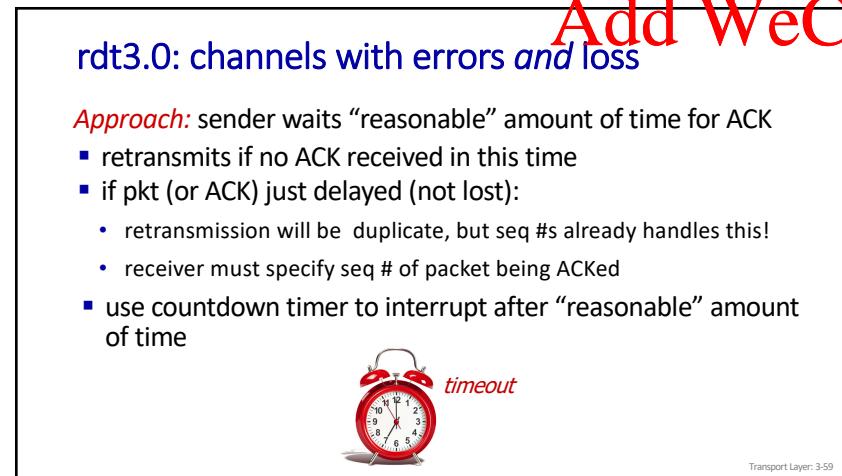


57

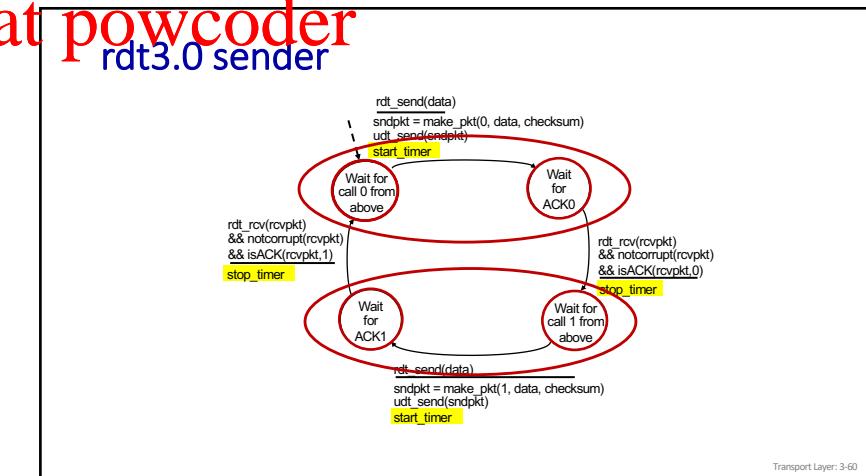


58

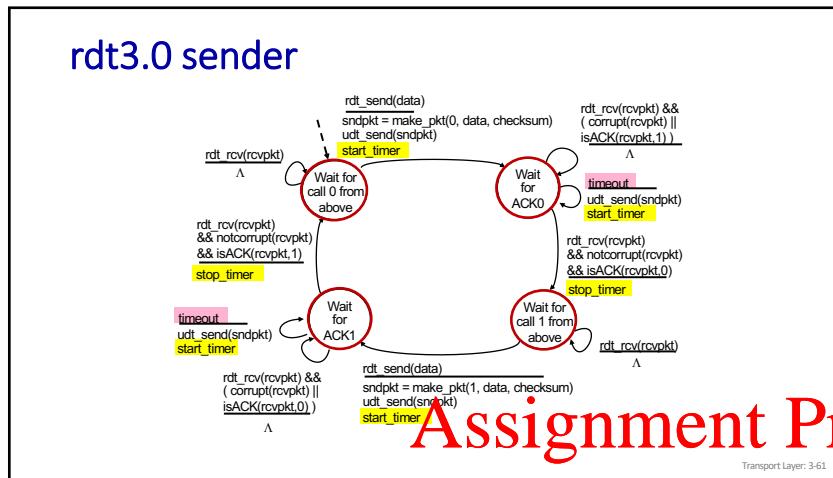
<https://powcoder.com>



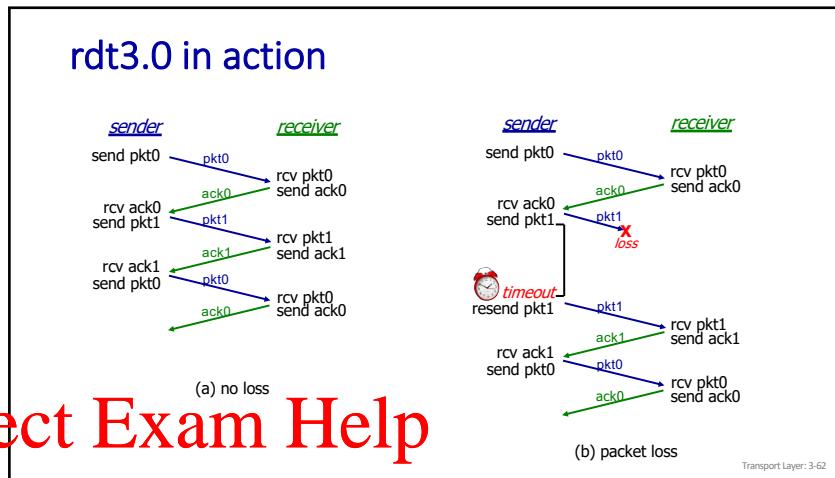
59



60

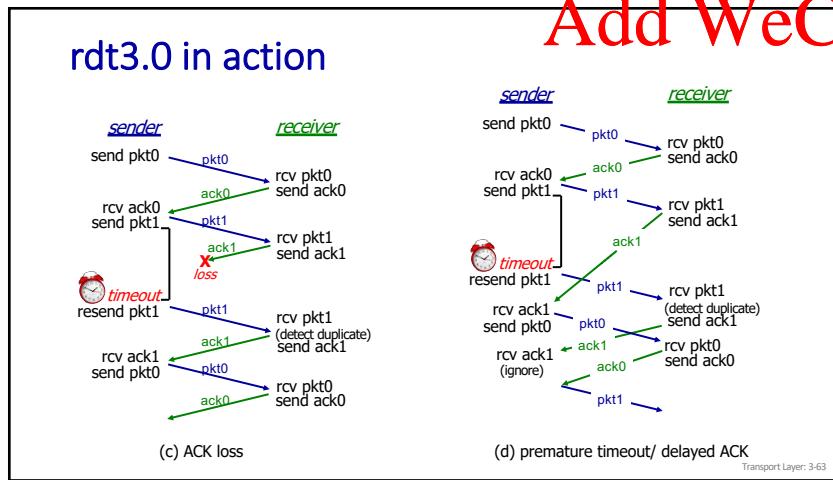


61

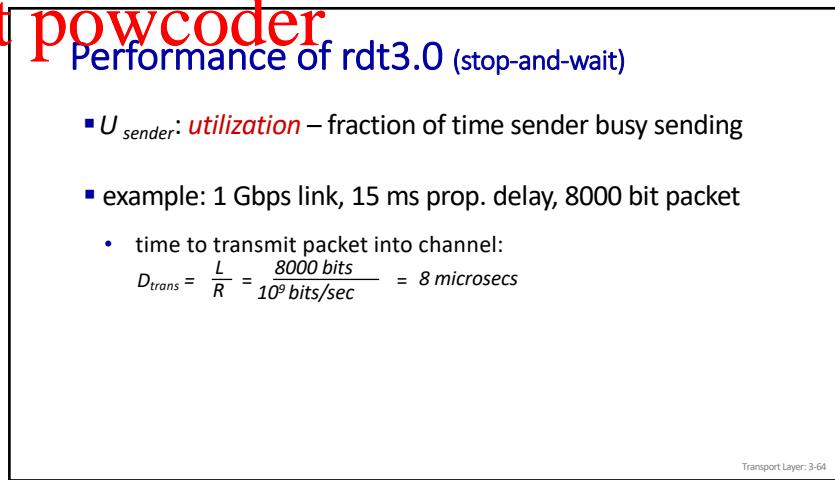


62

<https://powcoder.com>

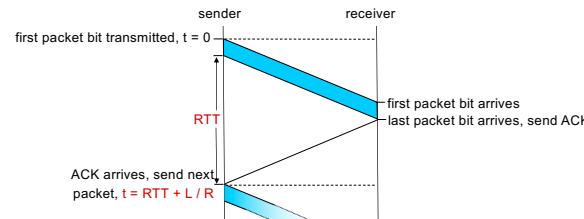


63



64

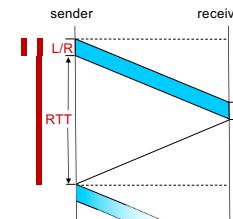
### rdt3.0: stop-and-wait operation



65

### rdt3.0: stop-and-wait operation

$$\begin{aligned} U_{\text{sender}} &= \frac{L / R}{RTT + L / R} \\ &= \frac{.008}{30.008} \\ &= 0.00027 \end{aligned}$$



- rdt 3.0 protocol performance stinks!
- Protocol limits performance of underlying infrastructure (channel)

Transport Layer: 3-66

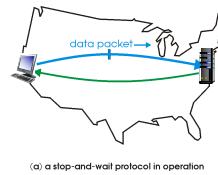
66

<https://powcoder.com>

### rdt3.0: pipelined protocols operation

**pipelining:** sender allows multiple, "in-flight", yet-to-be-acknowledged packets

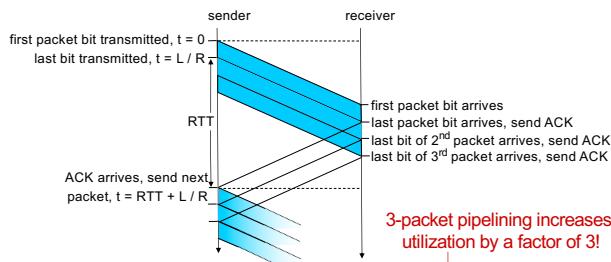
- range of sequence numbers must be increased
- buffering at sender and/or receiver



67

### Add WeChat powcoder

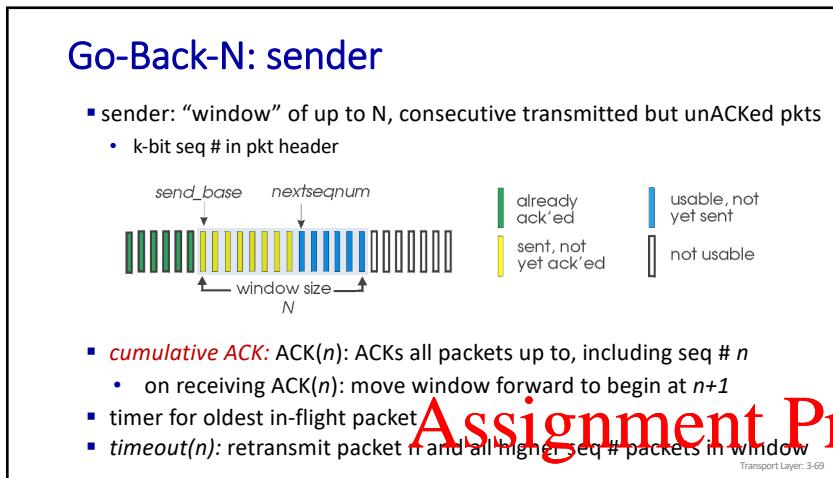
### Pipelining: increased utilization



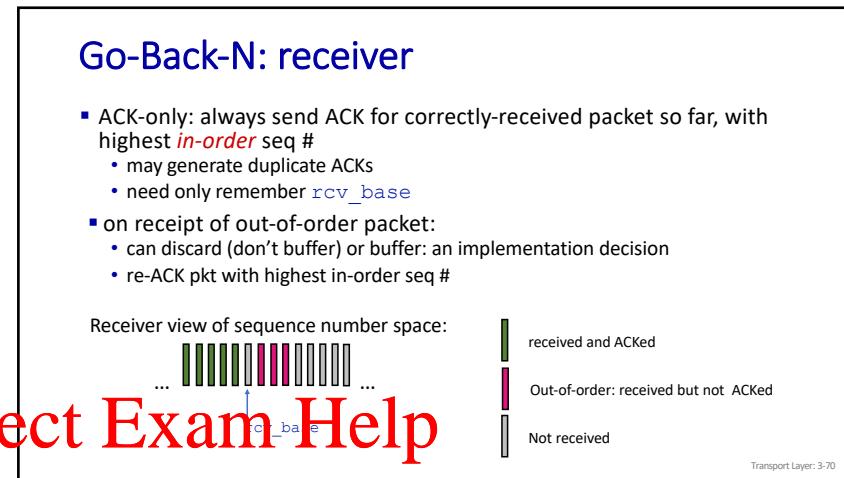
$$U_{\text{sender}} = \frac{3L / R}{RTT + L / R} = \frac{.0024}{30.008} = 0.00081$$

Transport Layer: 3-68

68

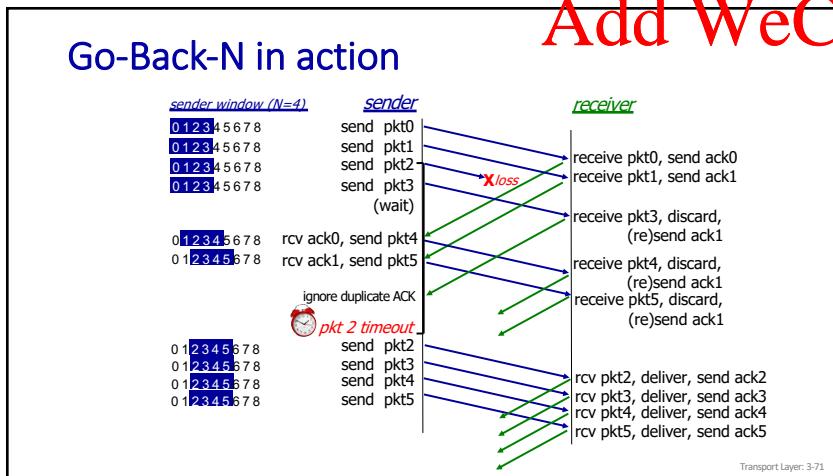


69

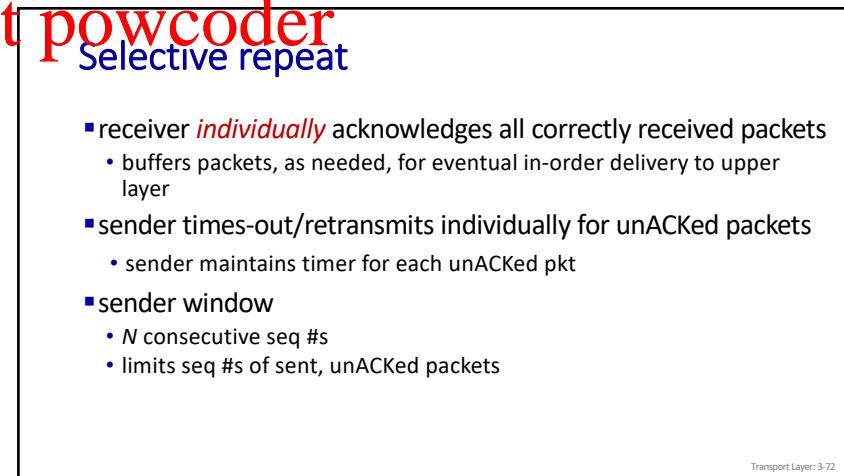


70

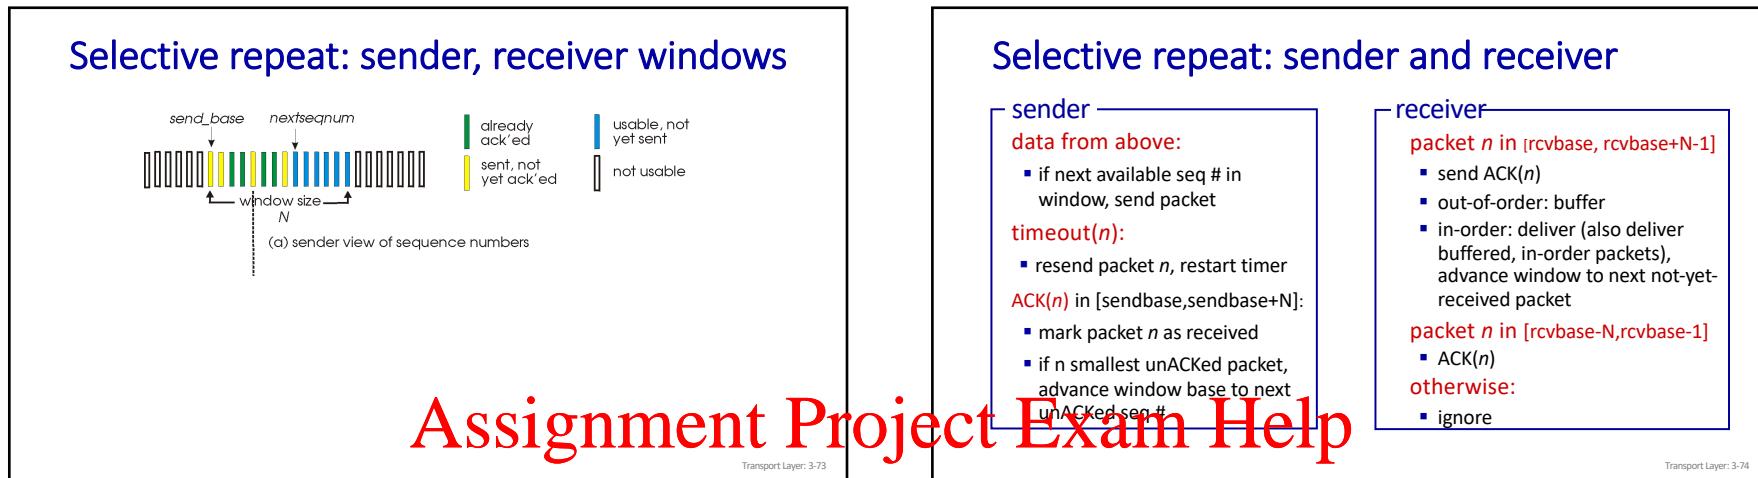
<https://powcoder.com>



71



72

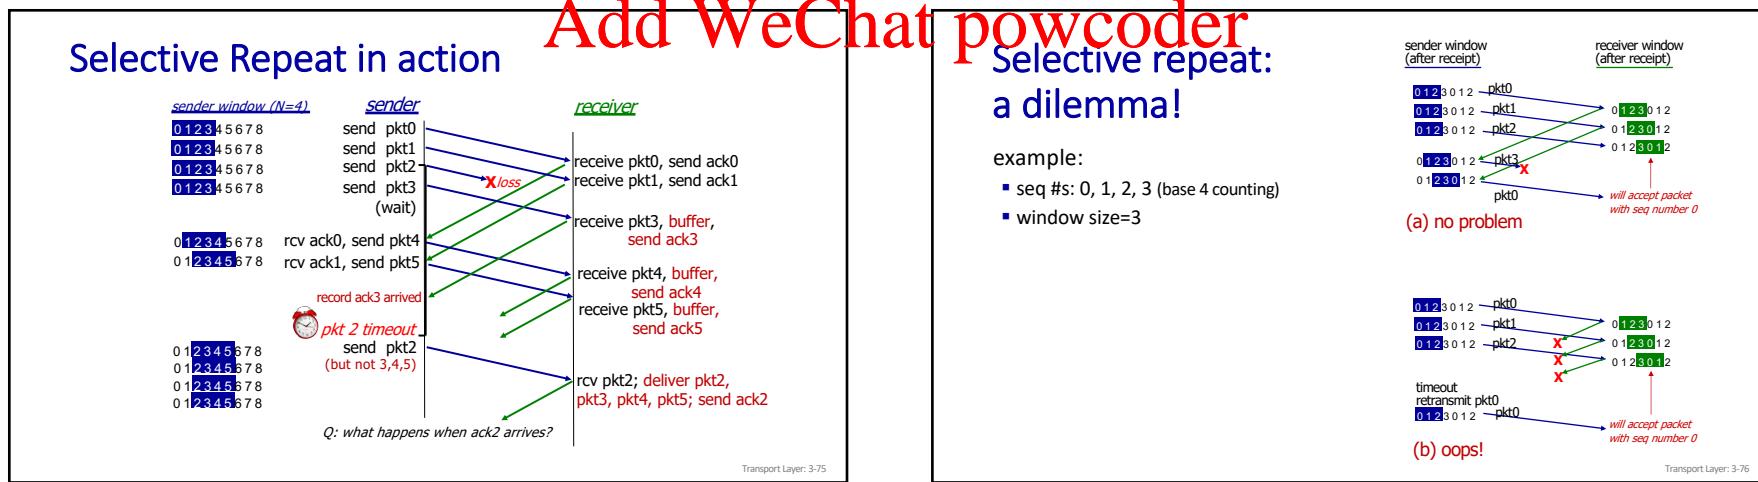


73

74

Assignment Project Exam Help

<https://powcoder.com>



75

76

**Selective repeat:  
a dilemma!**

example:

- seq #s: 0, 1, 2, 3 (base 4 counting)
- window size=3

Q: what relationship is needed between sequence # size and window size to avoid problem in scenario (b)?

77

## Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
- Principles of congestion control
- Congestion control



Transport Layer: 3-78

Assignment Project Exam Help

78

<https://powcoder.com>

**TCP: overview** RFCs: 793,1122, 2018, 5681, 7323

- point-to-point:**
  - one sender, one receiver
- reliable, in-order byte steam:**
  - no "message boundaries"
- full duplex data:**
  - bi-directional data flow in same connection
  - MSS: maximum segment size
- cumulative ACKs**
- pipelining:**
  - TCP congestion and flow control set window size
- connection-oriented:**
  - handshaking (exchange of control messages) initializes sender, receiver state before data exchange
- flow controlled:**
  - sender will not overwhelm receiver

79

**TCP segment structure**

source port #	dest port #
sequence number	acknowledgement number
length (of TCP header)	head not used
Internet checksum	C E R S F receive window
options (variable length)	
TCP options	
C, E: congestion notification	
RST, SYN, FIN: connection management	
application data (variable length)	

32 bits

segment seq #: counting bytes of data into bytestream (not segments!)

flow control: # bytes receiver willing to accept

data sent by application into TCP socket

80

## TCP sequence numbers, ACKs

**Sequence numbers:**

- byte stream “number” of first byte in segment’s data

**Acknowledgements:**

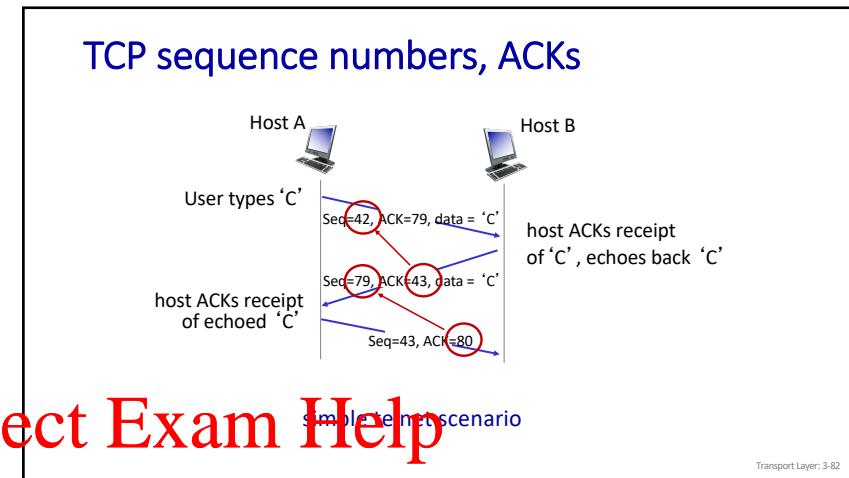
- seq # of next byte expected from other side
- cumulative ACK

**Q:** how receiver handles out-of-order segments

- A:** TCP spec doesn’t say, - up to implementor

Assignment Project Exam Help

81



82

https://powcoder.com

## TCP round trip time, timeout

**Q:** how to set TCP timeout value?

- longer than RTT, but RTT varies!
- too short:** premature timeout, unnecessary retransmissions
- too long:** slow reaction to segment loss

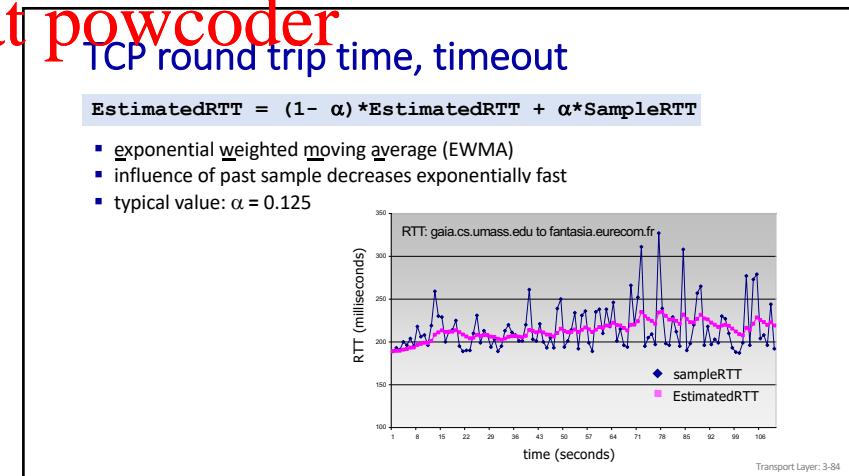
**Q:** how to estimate RTT?

- SampleRTT:** measured time from segment transmission until ACK receipt
  - ignore retransmissions
- SampleRTT** will vary, want estimated RTT “smoother”
  - average several *recent* measurements, not just current **SampleRTT**

Add WeChat powcoder

Transport Layer: 3-83

83



84

### TCP round trip time, timeout

- timeout interval: **EstimatedRTT** plus “safety margin”
  - large variation in **EstimatedRTT**: want a larger safety margin

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$

**DevRTT:** EWMA of **SampleRTT** deviation from **EstimatedRTT**:

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )

\* Check out the online interactive exercises for more examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)

Transport Layer: 3-85

### TCP Sender (simplified)

- event: data received from application**
  - create segment with seq #
  - seq # is byte-stream number of first data byte in segment
  - start timer if not already running
    - think of timer as for oldest unACKed segment
    - expiration interval: **TimeOutInterval**
- event: timeout**
  - retransmit segment that caused timeout
  - restart timer
- event: ACK received**
  - if ACK acknowledges previously unACKed segments
    - update what is known to be ACKed
  - start timer if there are still unACKed segments

Transport Layer: 3-86

85

86

Assignment Project Exam Help

<https://powcoder.com>

### TCP Receiver: ACK generation [RFC 5681]

Event at receiver	TCP receiver action

Transport Layer: 3-87

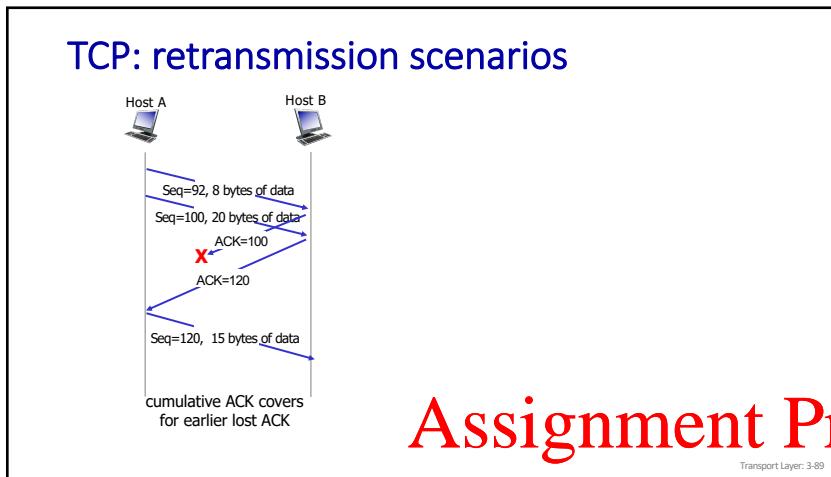
### Add WeChat powcoder

### TCP: retransmission scenarios

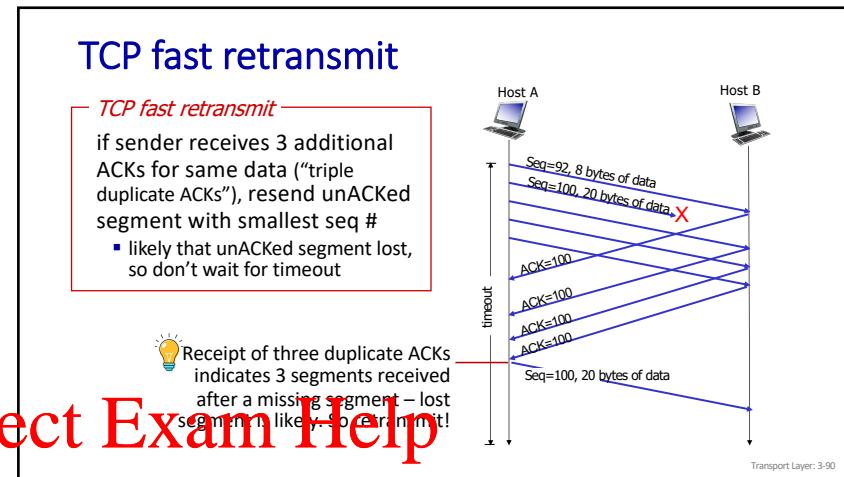
Transport Layer: 3-88

87

88



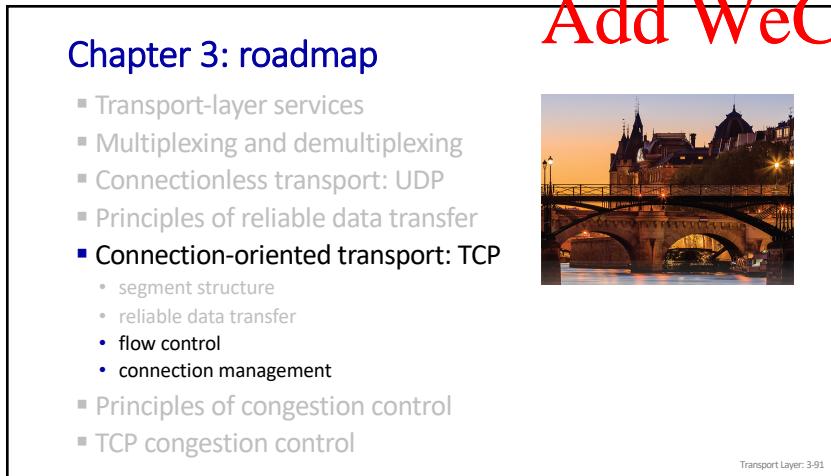
89



90

Assignment Project Exam Help

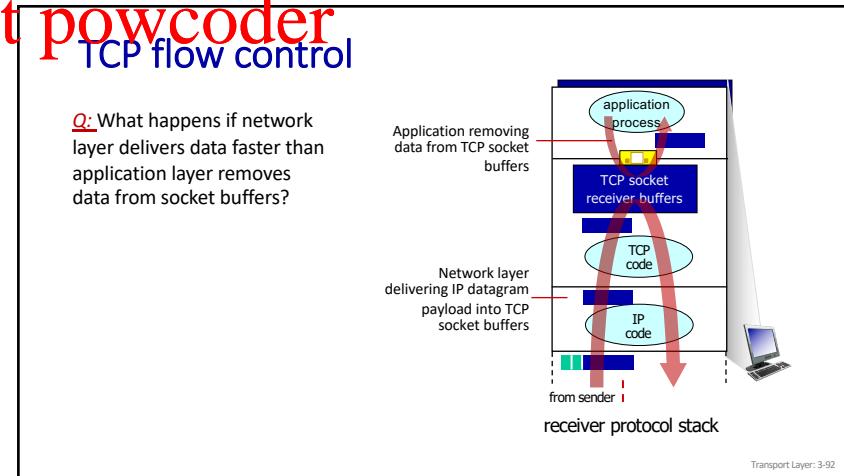
<https://powcoder.com>



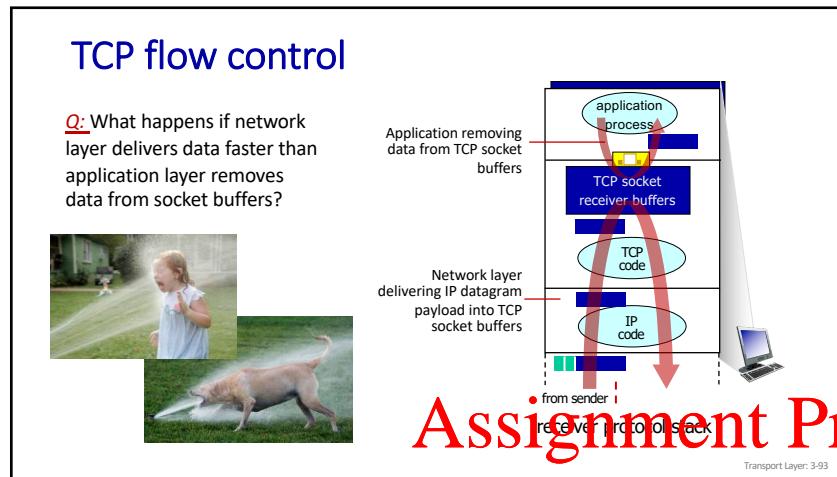
91

Add WeChat powcoder

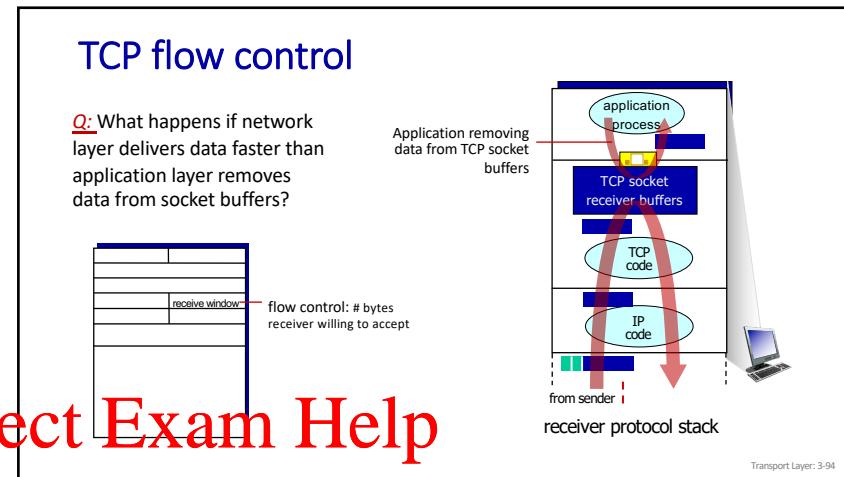
TCP flow control



92

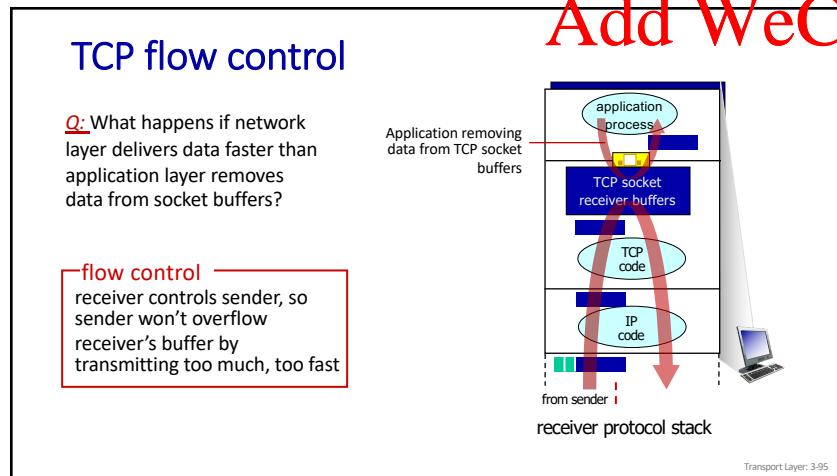


93

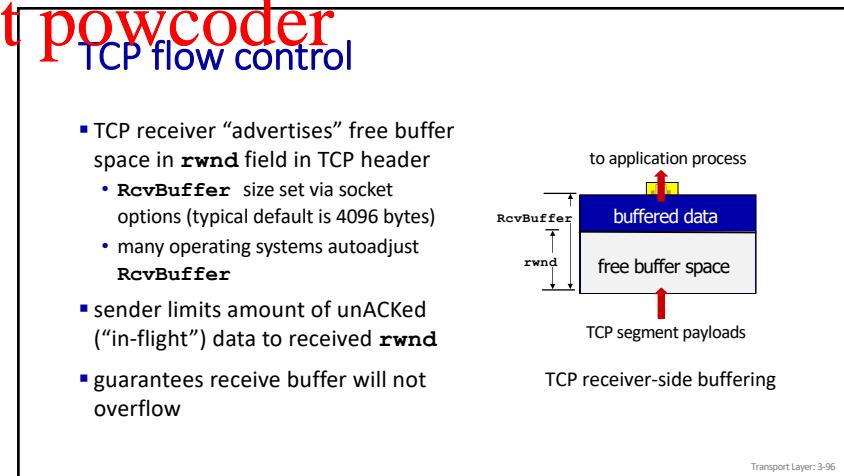


94

<https://powcoder.com>



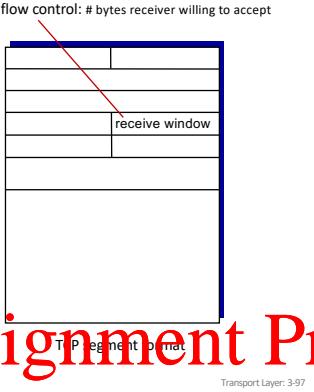
95



96

### TCP flow control

- TCP receiver “advertises” free buffer space in **rwnd** field in TCP header
  - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
  - many operating systems autoadjust **RcvBuffer**
- sender limits amount of unACKed (“in-flight”) data to received **rwnd**
- guarantees receive buffer will not overflow



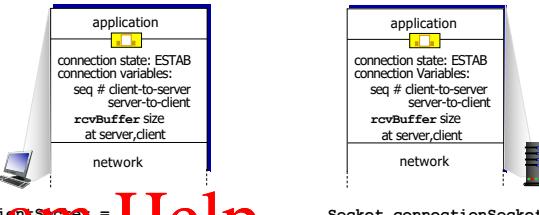
flow control: # bytes receiver willing to accept  
receive window

TCP segment format  
Transport Layer: 3-97

### TCP connection management

before exchanging data, sender/receiver “handshake”:

- agree to establish connection (each knowing the other willing to establish connection)
- agree on connection parameters (e.g., starting seq #s)



application  
connection state: ESTAB  
connection variables:  
seq # client-to-server  
server-to-client  
rcvBuffer size  
at server/client

application  
connection state: ESTAB  
connection Variables:  
seq # client-to-server  
server-to-client  
rcvBuffer size  
at server/client

Socket connectionSocket = welcomeSocket.accept();  
Transport Layer: 3-98

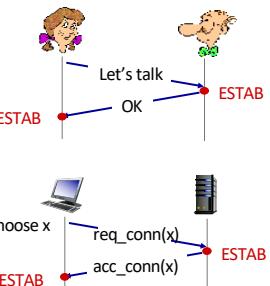
97

98

<https://powcoder.com>

### Agreeing to establish a connection

2-way handshake:

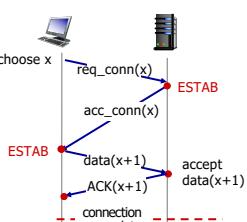


Let's talk  
ESTAB  
OK  
ESTAB

choose x  
req\_conn(x)  
ESTAB  
ESTAB  
acc\_conn(x)

Transport Layer: 3-99

### 2-way handshake scenarios



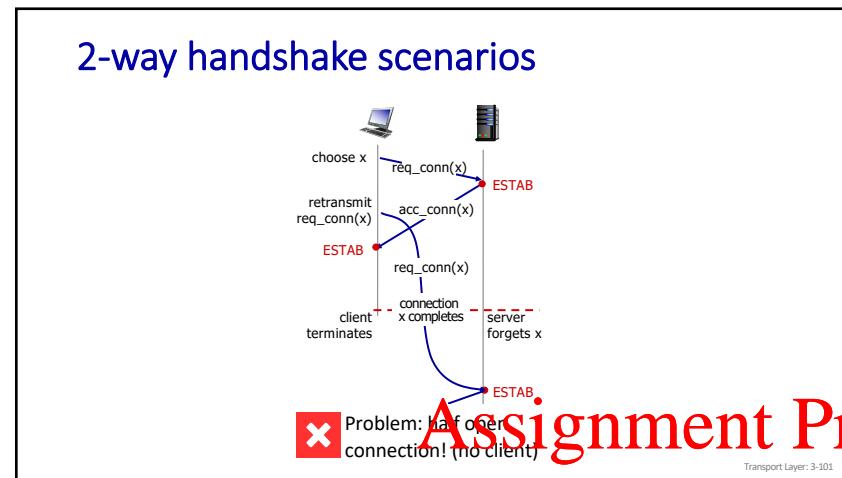
choose x  
req\_conn(x)  
ESTAB  
ESTAB  
acc\_conn(x)  
data(x+1)  
ACK(x+1)  
accept data(x+1)  
connection x completes

No problem! 

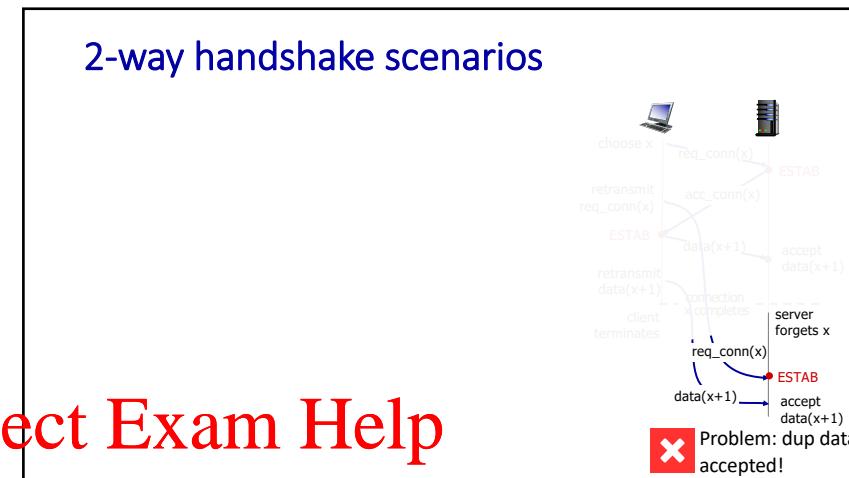
Transport Layer: 3-100

99

100

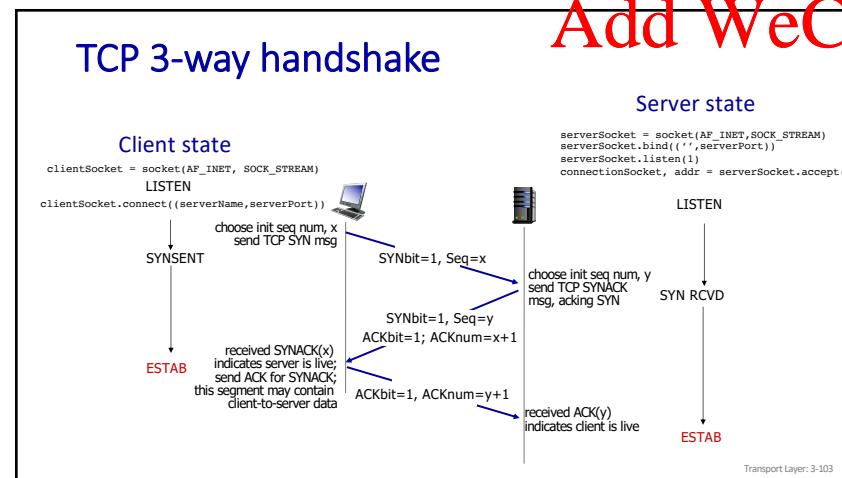


101

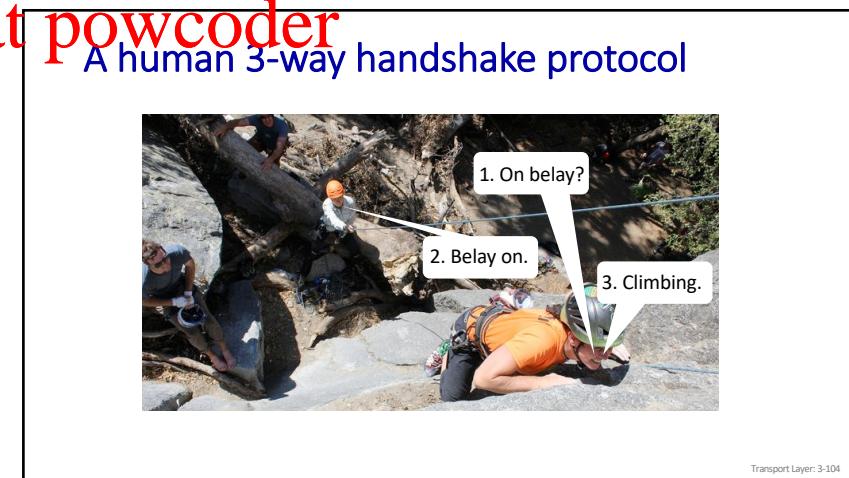


102

<https://powcoder.com>



103



104

## Closing a TCP connection

- client, server each close their side of connection
  - send TCP segment with FIN bit = 1
- respond to received FIN with ACK
  - on receiving FIN, ACK can be combined with own FIN
- simultaneous FIN exchanges can be handled

105

Transport Layer: 3-105

# Assignment Project Exam Help



Transport Layer: 3-106

## Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer services

106

<https://powcoder.com>

## Principles of congestion control

### Congestion:

- informally: “too many sources sending too much data too fast for *network* to handle”
- manifestations:
  - long delays (queueing in router buffers)
  - packet loss (buffer overflow at routers)
- different from flow control!
- a top-10 problem!



**congestion control:**  
too many senders,  
sending too fast

**flow control:** one sender  
too fast for one receiver

Transport Layer: 3-107

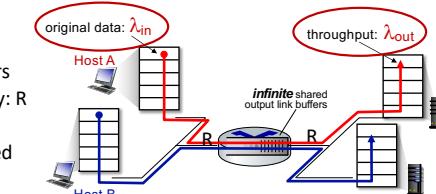
107

## Add WeChat powcoder

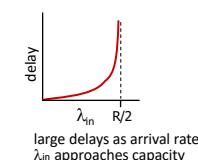
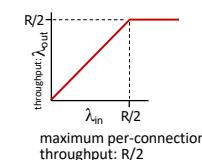
### Causes/costs of congestion: scenario 1

#### Simplest scenario:

- one router, infinite buffers
- input, output link capacity: R
- two flows
- no retransmissions needed

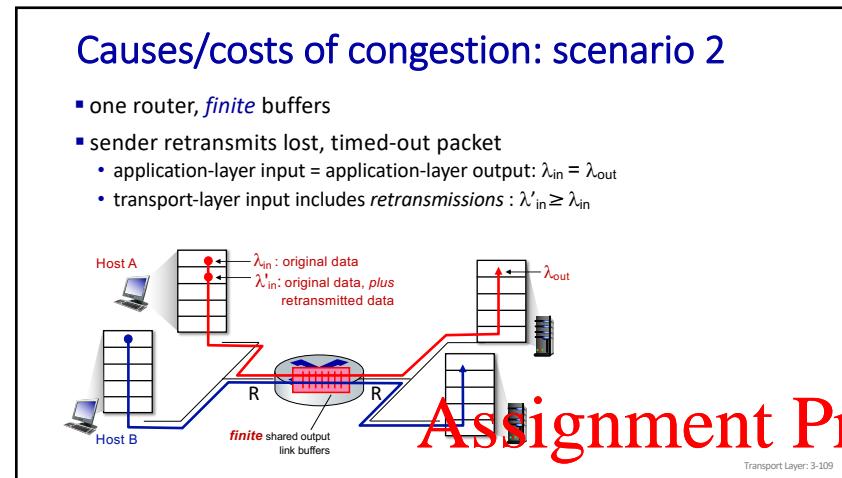


**Q:** What happens as  
arrival rate  $\lambda_{in}$   
approaches  $R/2$ ?

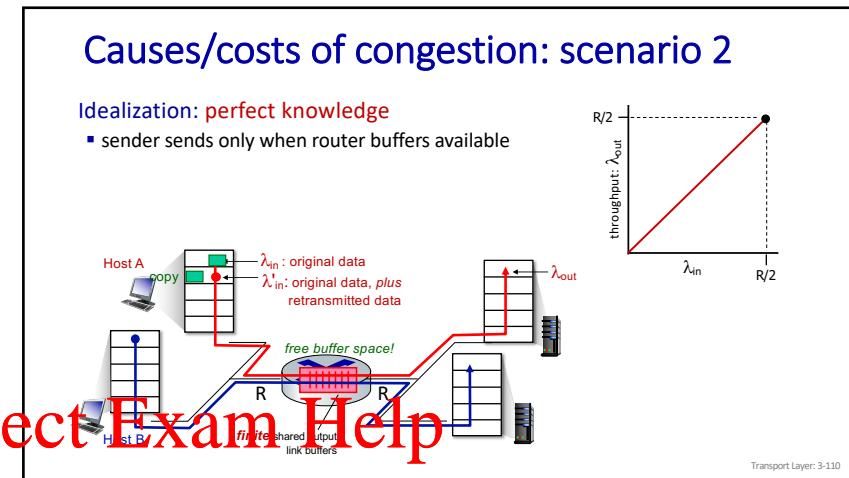


Transport Layer: 3-108

108

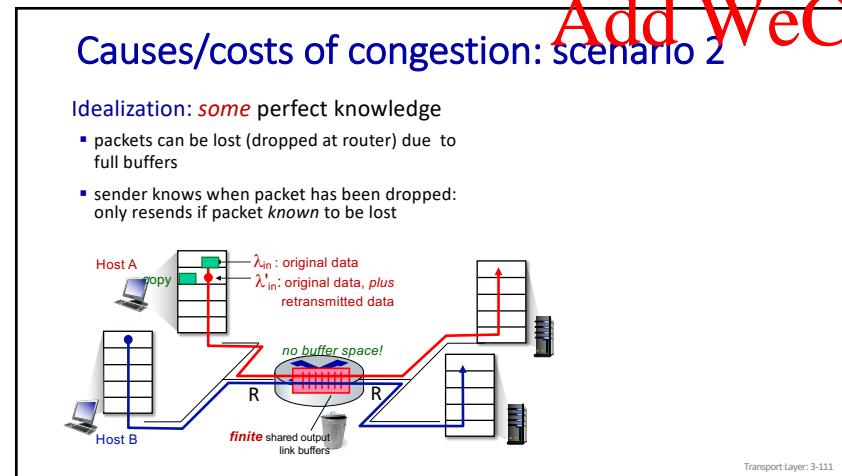


109

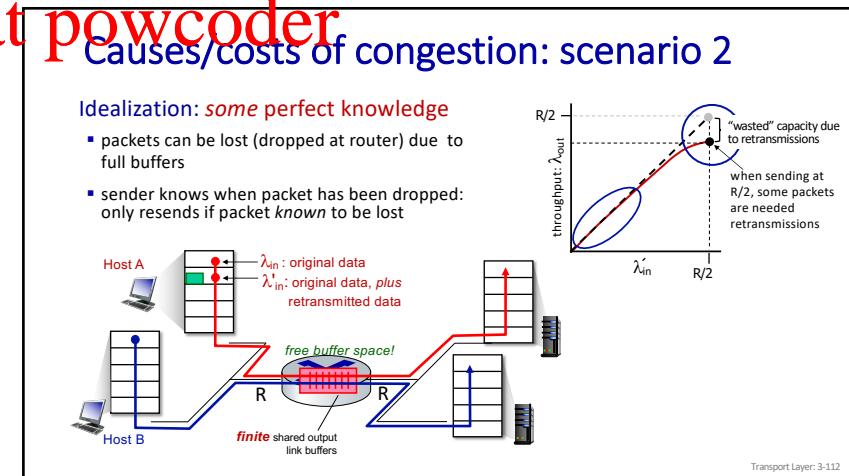


110

Assignment Project Exam Help  
https://powcoder.com



111



112

### Causes/costs of congestion: scenario 2

**Realistic scenario: un-needed duplicates**

- packets can be lost, dropped at router due to full buffers – requiring retransmissions
- but sender times can time out prematurely, sending **two** copies, **both** of which are delivered

Host A  
Host B  
Router R  
 $\lambda'_{in}$ : original data  
 $\lambda'_{in}$ : original data, plus retransmitted data  
finite shared output link buffers  
free buffer space!

throughput:  $\lambda'_{out}$

"wasted" capacity due to un-needed retransmissions

when sending at  $R/2$ , some packets are retransmissions, including needed and un-needed duplicates, that are delivered!

Transport Layer: 3-113

113

### Causes/costs of congestion: scenario 2

**Realistic scenario: un-needed duplicates**

- packets can be lost, dropped at router due to full buffers – requiring retransmissions
- but sender times can time out prematurely, sending **two** copies, **both** of which are delivered

Host A  
Host B  
Router R  
 $\lambda'_{in}$ : original data  
 $\lambda'_{in}$ : original data, plus retransmitted data  
finite shared output link buffers  
free buffer space!

throughput:  $\lambda'_{out}$

"wasted" capacity due to un-needed retransmissions

when sending at  $R/2$ , some packets are retransmissions, including needed and un-needed duplicates, that are delivered!

Transport Layer: 3-114

114

Assignment Project Exam Help

<https://powcoder.com>

### Causes/costs of congestion: scenario 3

**four senders**

**multi-hop paths**

**timeout/retransmit**

**Q:** what happens as  $\lambda_{in}$  and  $\lambda'_{in}$  increase?

**A:** as red  $\lambda'_{in}$  increases, all arriving blue pkts at upper queue are dropped, blue throughput  $\rightarrow 0$

Host A  
Host B  
Host C  
Host D  
 $\lambda_{in}$ : original data  
 $\lambda'_{in}$ : original data, plus retransmitted data  
finite shared output link buffers  
free buffer space!

throughput:  $\lambda'_{out}$

Transport Layer: 3-115

115

### Causes/costs of congestion: scenario 3

**another “cost” of congestion:**

- when packet dropped, any upstream transmission capacity and buffering used for that packet was wasted!

Host A  
Host B  
Host C  
Host D  
 $\lambda_{in}$ : original data  
 $\lambda'_{in}$ : original data, plus retransmitted data  
finite shared output link buffers  
free buffer space!

throughput:  $\lambda'_{out}$

Transport Layer: 3-116

116

## Causes/costs of congestion: insights

- throughput can never exceed capacity
- delay increases as capacity approached
- loss/retransmission decreases effective throughput
- un-needed duplicates further decreases effective throughput
- upstream transmission capacity / buffering wasted for packets lost downstream

Transport Layer: 3-117

117

## Approaches towards congestion control

### End-end congestion control:

- no explicit feedback from network
- congestion *inferred* from observed loss, delay
- approach taken by TCP

Transport Layer: 3-118

118

Assignment Project Exam Help  
<https://powcoder.com>

## Approaches towards congestion control

### Network-assisted congestion control:

- routers provide *direct* feedback to sending/receiving hosts with flows passing through congested router
- may indicate congestion level or explicitly set sending rate
- TCP ECN, ATM, DECBIT protocols

Transport Layer: 3-119

119

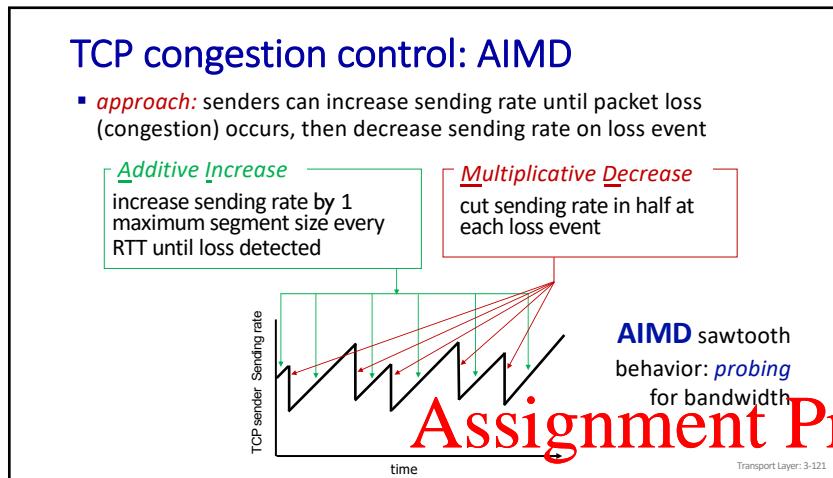
## Add WeChat powcoder

### Chapter 3: roadmap

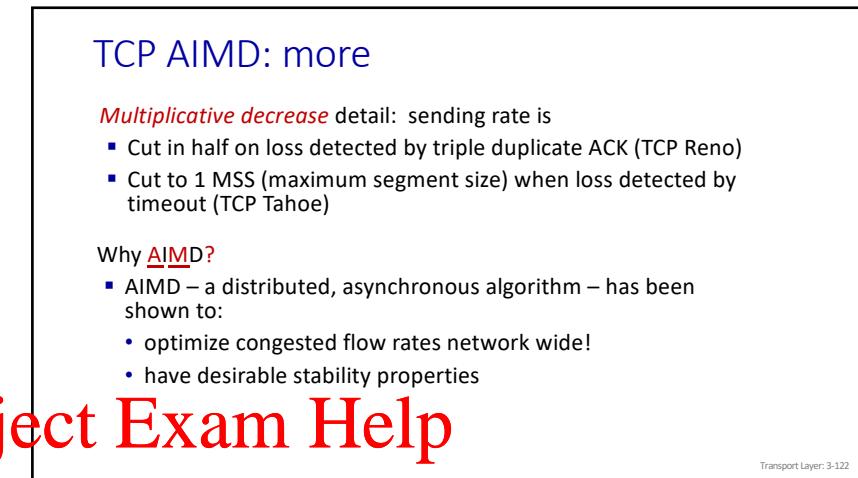
- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer functionality

Transport Layer: 3-120

120

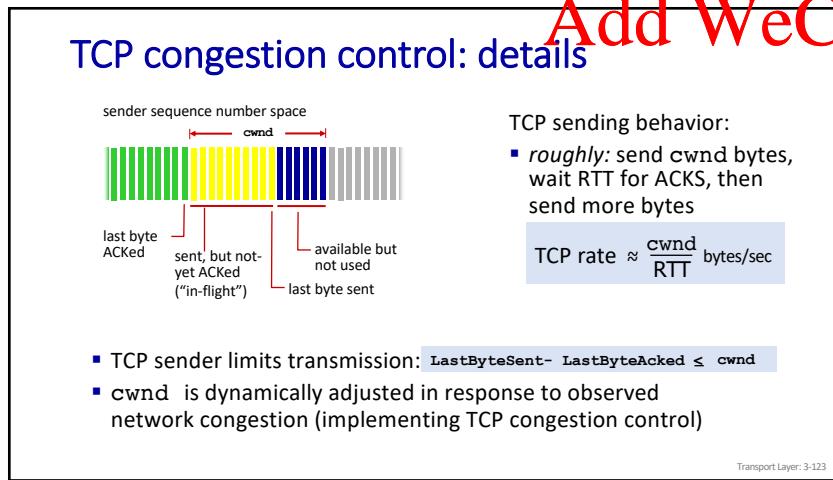


121

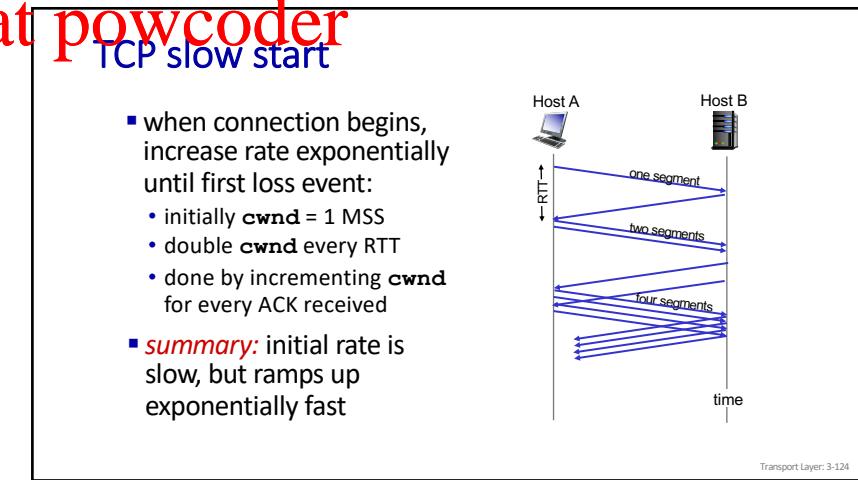


122

<https://powcoder.com>



123



124

## TCP: from slow start to congestion avoidance

**Q:** when should the exponential increase switch to linear?

**A:** when **cwnd** gets to 1/2 of its value before timeout.

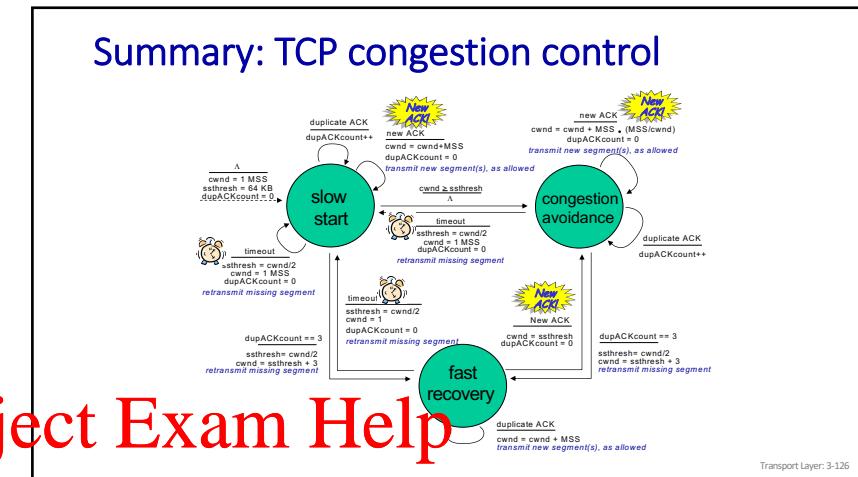
**Implementation:**

- variable **ssthresh**
- on loss event, **ssthresh** is set to 1/2 of **cwnd** just before loss event

\* Check out the online interactive exercises for more examples: <http://cse.csail.mit.edu/6.S095/exercises/>

Transport Layer: 3-125

125



126

# Assignment Project Exam Help

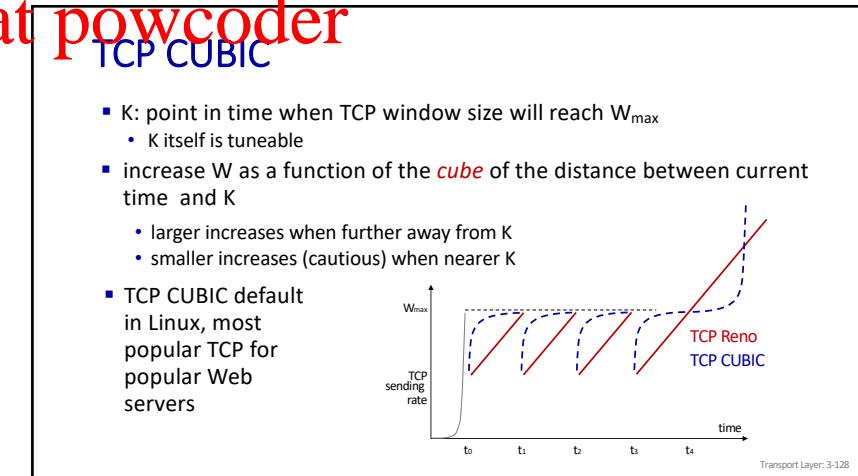
<https://powcoder.com>

## TCP CUBIC

- Is there a better way than AIMD to “probe” for usable bandwidth?
- Insight/intuition:
  - $W_{\max}$ : sending rate at which congestion loss was detected
  - congestion state of bottleneck link probably (?) hasn't changed much
  - after cutting rate/window in half on loss, initially ramp to  $W_{\max}$  *faster*, but then approach  $W_{\max}$  more *slowly*

Transport Layer: 3-127

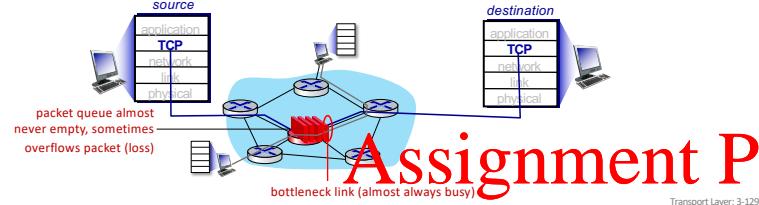
127



128

## TCP and the congested “bottleneck link”

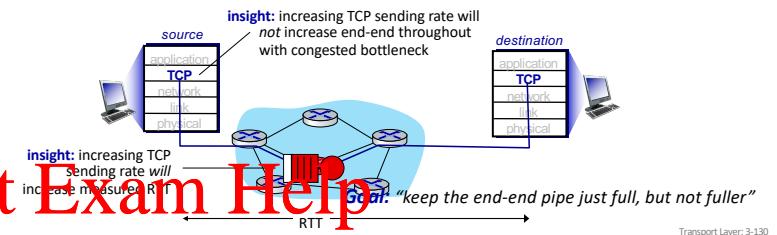
- TCP (classic, CUBIC) increase TCP's sending rate until packet loss occurs at some router's output: the *bottleneck link*



129

## TCP and the congested “bottleneck link”

- TCP (classic, CUBIC) increase TCP's sending rate until packet loss occurs at some router's output: the *bottleneck link*
- understanding congestion: useful to focus on congested bottleneck link



130

<https://powcoder.com>

## Delay-based TCP congestion control

Keeping sender-to-receiver pipe “just full enough, but no fuller”: keep bottleneck link busy transmitting, but avoid high delays/buffering



### Delay-based approach:

- RTT<sub>min</sub> - minimum observed RTT (uncongested path)
- uncongested throughput with congestion window cwnd<sub>0</sub> is cwnd<sub>0</sub>/RTT<sub>min</sub>
- if measured throughput “very close” to uncongested throughput
  - increase cwnd linearly /\* since path not congested \*/
- else if measured throughput “far below” uncongested throughput
  - decrease cwnd linearly /\* since path is congested \*/

Transport Layer: 3-131

131

## Add WeChat powcoder

## Delay-based TCP congestion control

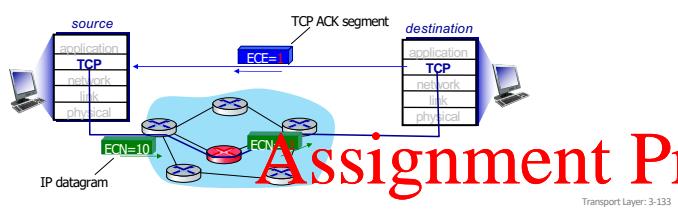
- congestion control without inducing/forcing loss
- maximizing throughout (“keeping the just pipe full... ”) while keeping delay low (“...but not fuller”)
- a number of deployed TCPs take a delay-based approach
  - BBR deployed on Google’s (internal) backbone network

Transport Layer: 3-132

132

## Explicit congestion notification (ECN)

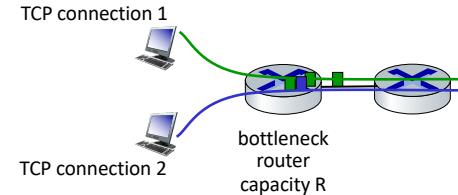
- TCP deployments often implement *network-assisted* congestion control:
- two bits in IP header (ToS field) marked *by network router* to indicate congestion
    - policy* to determine marking chosen by network operator
  - congestion indication carried to destination
  - destination sets ECE bit on ACK segment to notify sender of congestion
  - involves both IP (IP header ECN bit marking) and TCP (TCP header C,E bit marking)



133

## TCP fairness

**Fairness goal:** if  $K$  TCP sessions share same bottleneck link of bandwidth  $R$ , each should have average rate of  $R/K$



Transport Layer: 3-134

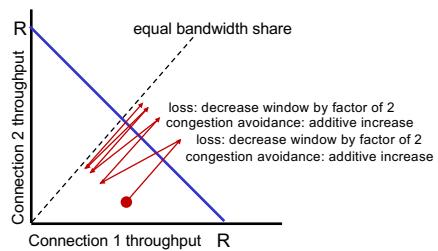
134

Assignment Project Exam Help  
<https://powcoder.com>

## Q: is TCP Fair?

Example: two competing TCP sessions:

- additive increase gives slope of 1, as throughput increases
- multiplicative decrease decreases throughput proportionally



Transport Layer: 3-135

135

Add WeChat powcoder

**Fairness:** must all network apps be “fair”?

### Fairness and UDP

- multimedia apps often do not use TCP
  - do not want rate throttled by congestion control
- instead use UDP:
  - send audio/video at constant rate, tolerate packet loss
- there is no “Internet police” policing use of congestion control

### Fairness, parallel TCP connections

- application can open *multiple* parallel connections between two hosts
- web browsers do this, e.g., link of rate  $R$  with 9 existing connections:
  - new app asks for 1 TCP, gets rate  $R/10$
  - new app asks for 11 TCPs, gets  $R/2$

Transport Layer: 3-136

136

### Transport layer: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- Principles of reliable data transfer
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer functionality



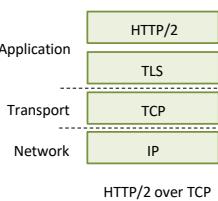
Transport Layer: 3-137

# Assignment Project Exam Help

137
138

### QUIC: Quick UDP Internet Connections

- application-layer protocol, on top of UDP
  - increase performance of HTTP
  - deployed on many Google servers, apps (Chrome, mobile YouTube app)



Transport Layer: 3-139

### Add WeChat powcoder

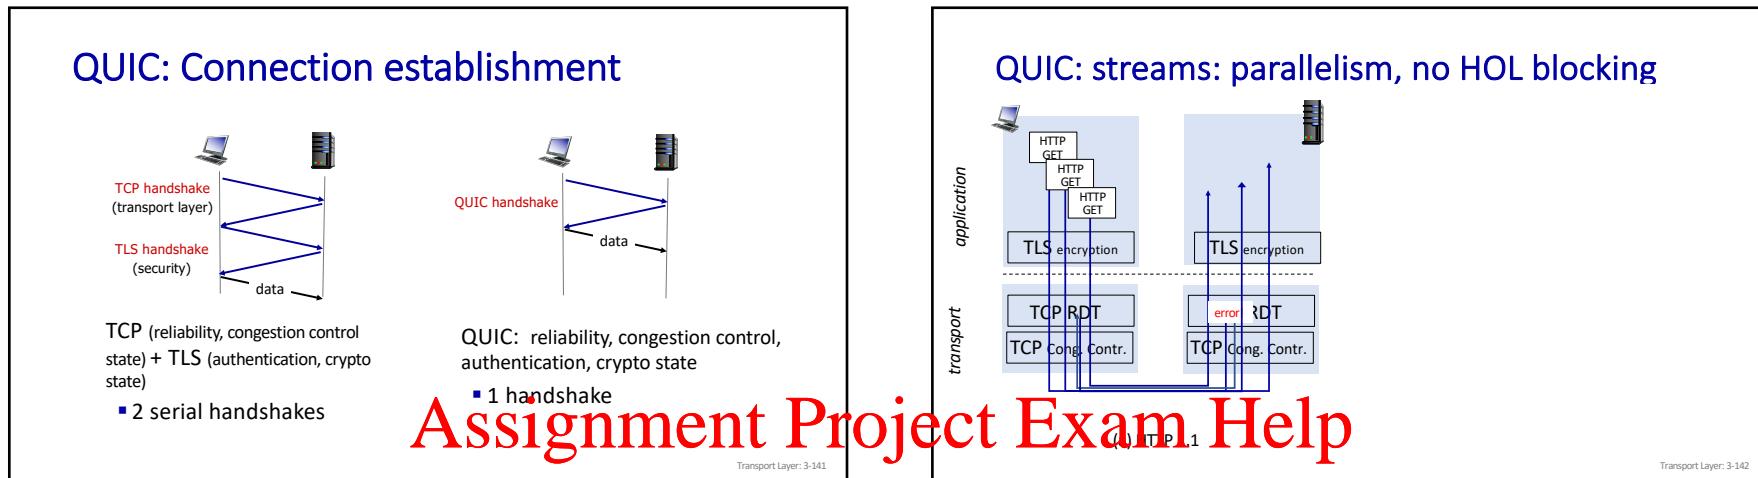
### QUIC: Quick UDP Internet Connections

adopts approaches we've studied in this chapter for connection establishment, error control, congestion control

- **error and congestion control:** “Readers familiar with TCP’s loss detection and congestion control will find algorithms here that parallel well-known TCP ones.” [from QUIC specification]
- **connection establishment:** reliability, congestion control, authentication, encryption, state established in one RTT

- multiple application-level “streams” multiplexed over single QUIC connection
  - separate reliable data transfer, security
  - common congestion control

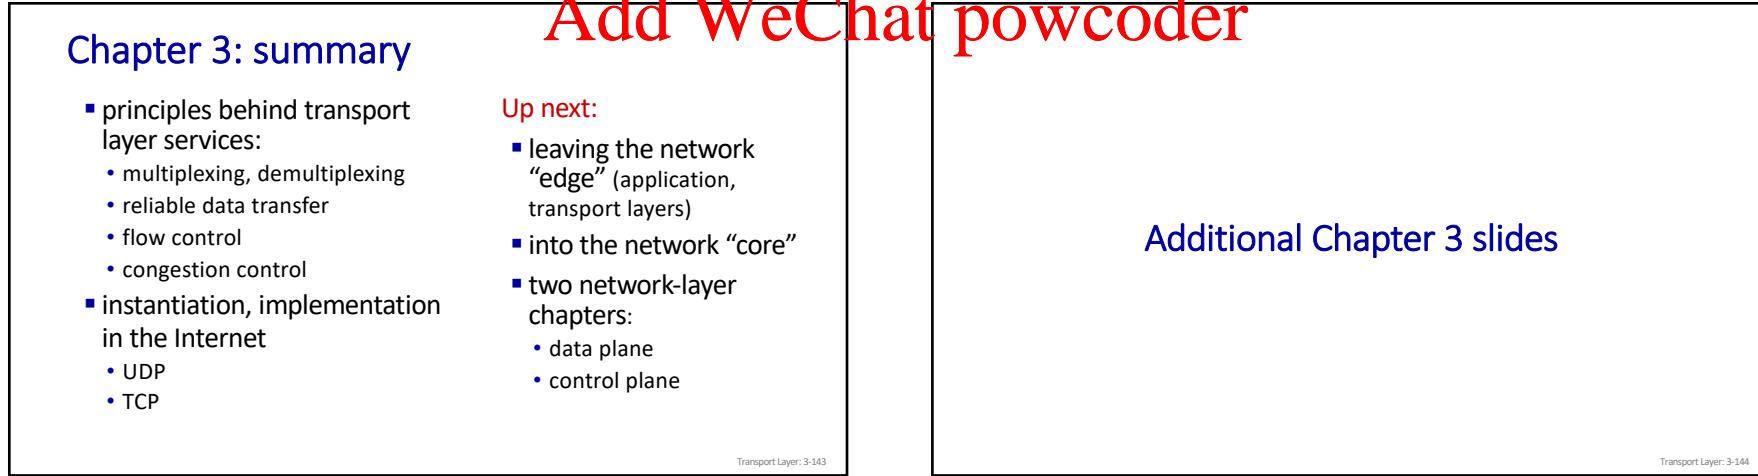
139
140



141

142

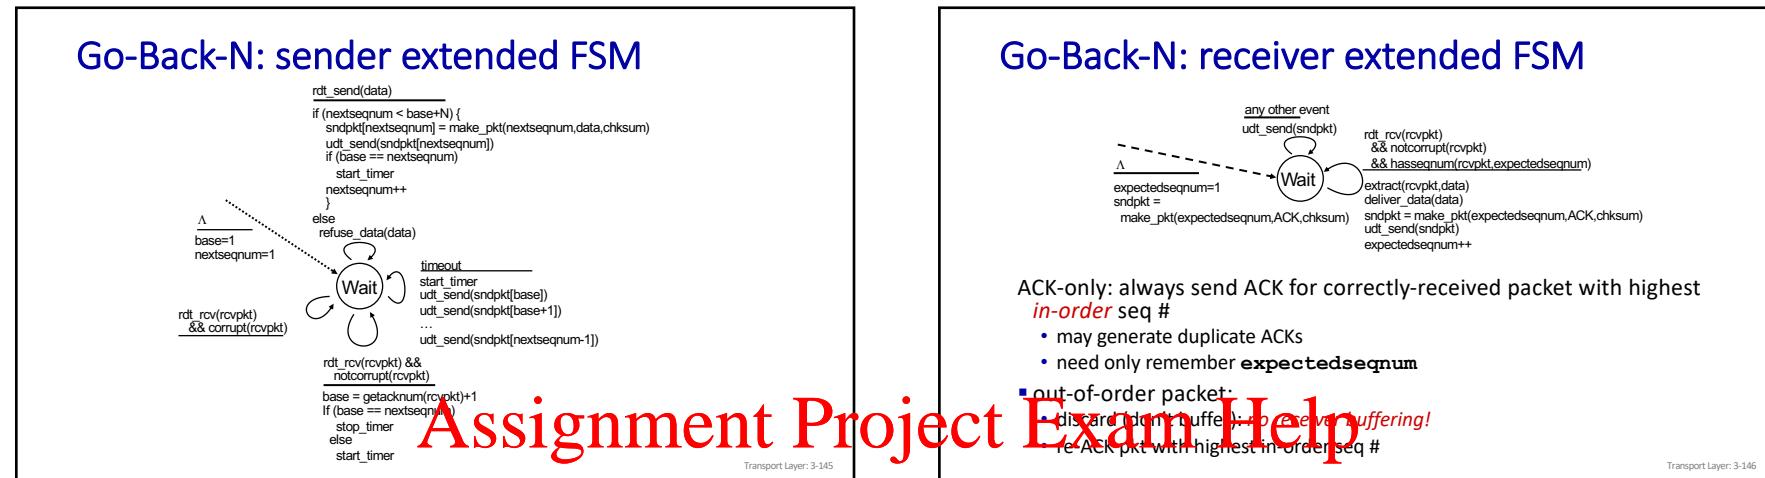
<https://powcoder.com>



143

144

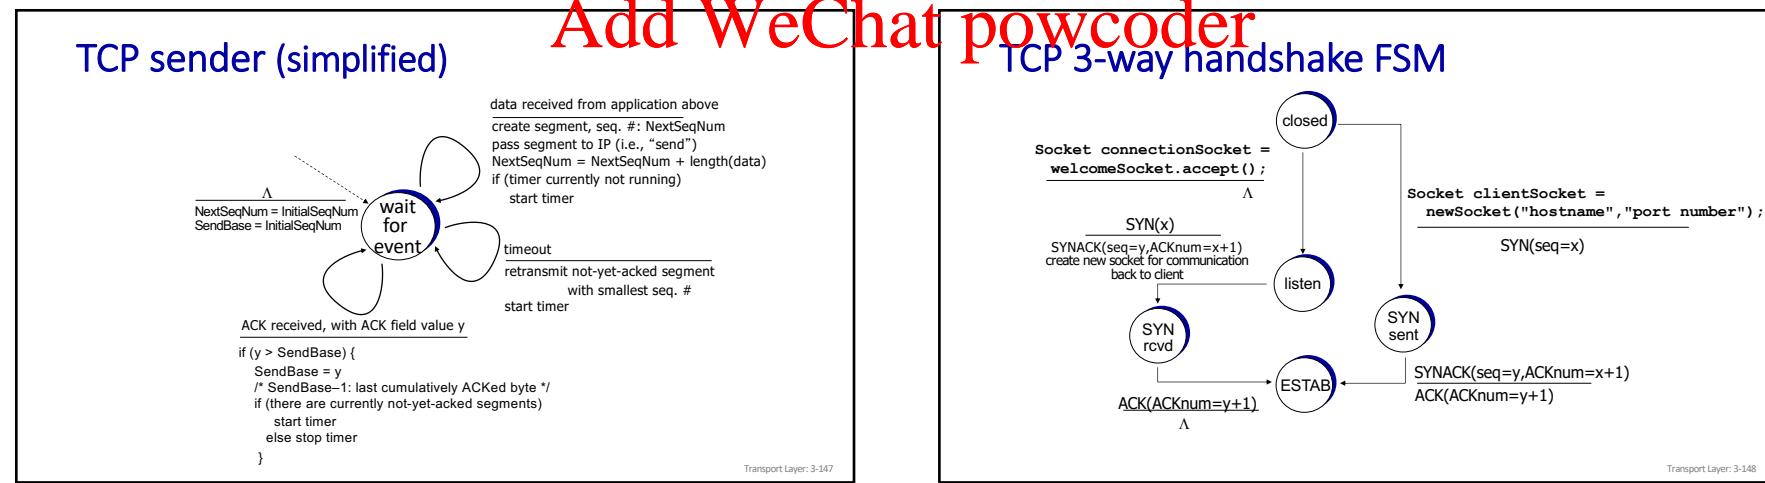
**Additional Chapter 3 slides**

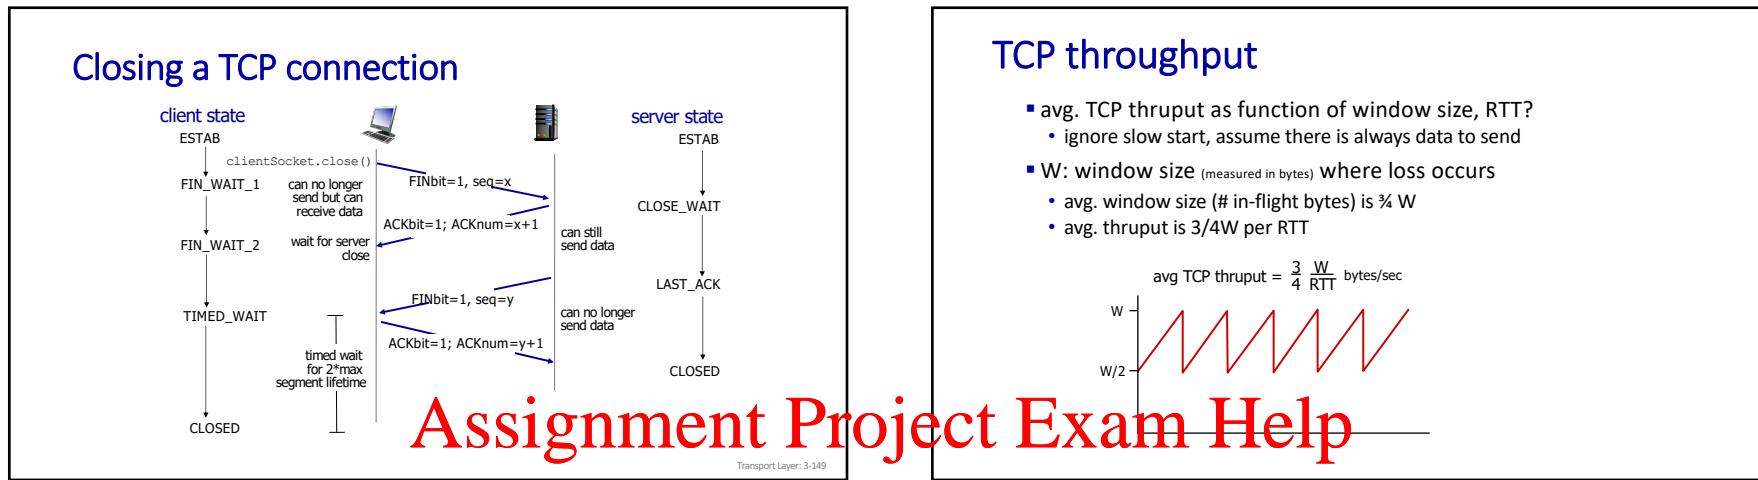


145

146

<https://powcoder.com>

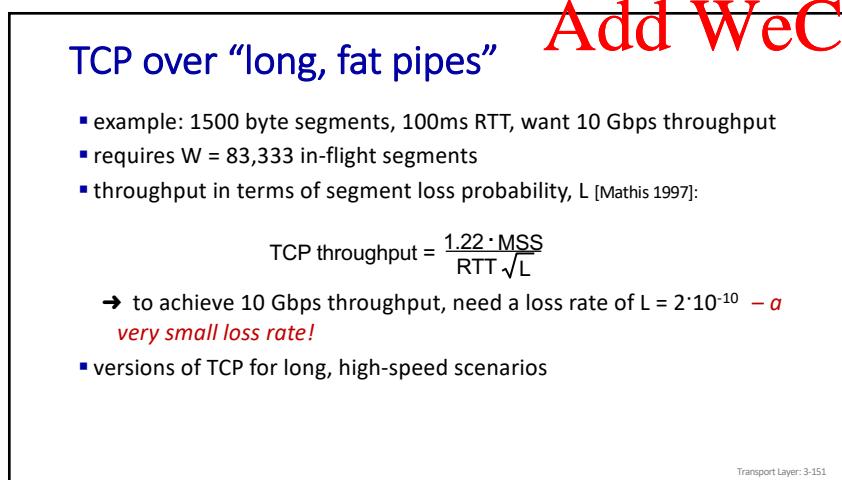




149

150

# Assignment Project Exam Help



151